

Multiprocessors

Feb-Jun 2021

4

Open Multi-Processing


Alejandro Guajardo Moreno

Recap... OpenMP Synchronization

Synchronization - bringing two or more threads to a known and well defined point in their execution.

The 2 most used flavors of synchronization:

➤ Barrier




➤ Mutual Exclusion

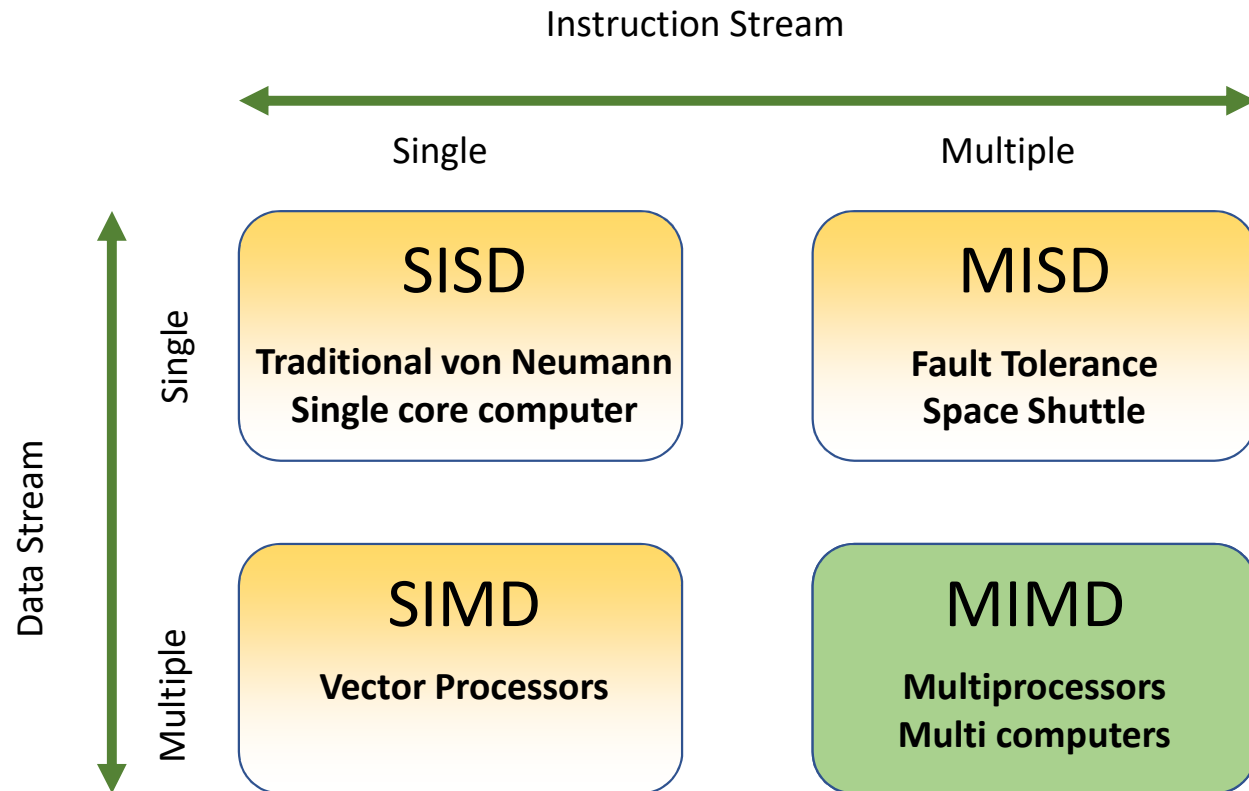
➤ Critical
➤ Atomic



Synchronization is **EXPENSIVE!**



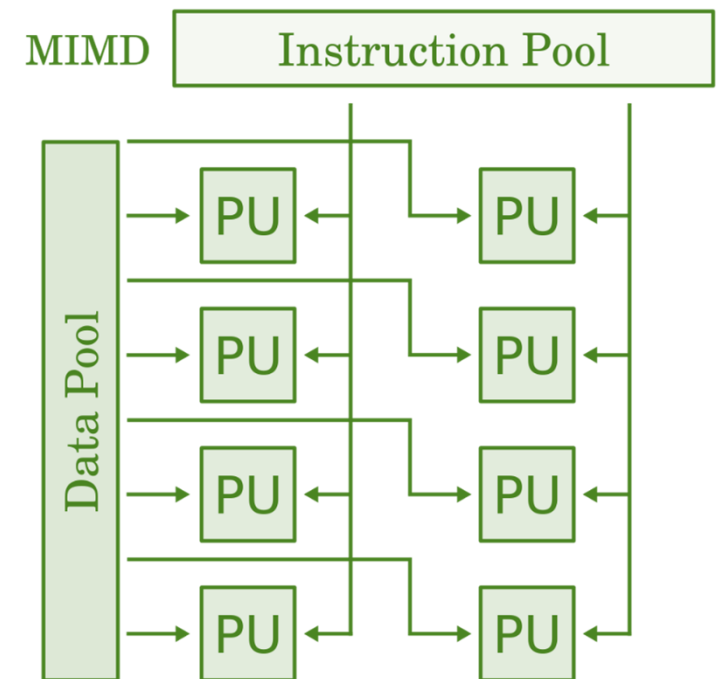
Flynn's Taxonomy



OpenMP Foucs: MIMD

SPMD (MIMD)

- **SPMD (Single Program, Multiple Data)**
- It is a subcategory of MIMD.
- Is the most common style of parallel programming.
- Tasks are split up and run **simultaneously** on **multiple** processors.
- Multiple autonomous processors simultaneously execute the same program at **independent** points
- Use general purpose CPUs.





OpenMP - Worksharing



Single construct shared
between the threads.

Worksharing

- Loop Construct
- Sections/Section Constructs
- Single Construct
- Task Construct

OpenMP – Worksharing – LOOP CONSTRUCT

```
for (i = 0; i < SIZE; i++)  
    A[i] += B[i];
```



Until today, we would
write it this way:

```
#pragma omp parallel  
{  
    int threadID, threadAmount, i, threadStart, threadEnd;  
    threadID = omp_get_thread_num();  
    threadAmount = omp_get_num_threads();  
    threadStart = threadID * SIZE / threadAmount;  
    threadEnd = (threadID + 1) * SIZE / threadAmount;  
    if (threadID == threadAmount - 1) threadEnd = SIZE;  
    for (i = threadStart; i < threadEnd; i++)  
        A[i] += B[i];  
}
```

Worksharing

Loop Construct

Sections/Section Constructs

Single Construct

Task Construct

OpenMP – Worksharing – LOOP CONSTRUCT



- I want to split iterations of the **for** loop to the different threads.
- I don't want to manually distribute the iterations.

Worksharing

Loop Construct

Sections/Section Constructs
Single Construct
Task Construct

Construct: `#pragma omp for`

```
for (i = 0; i < SIZE; i++)  
    A[i] += B[i];
```



```
#pragma omp parallel  
{  
    #pragma omp for  
    for (i = 0; i < SIZE; i++)  
        A[i] += B[i];  
}
```

❖ In OpenMP, the loop control index on a parallel loop is private to the thread.

OpenMP – Loop carried dependencies

- ✓ Happens when an iteration of a loop uses a value from a previous iteration.

```
int x,y;  
float A[SIZE];  
y = 10;  
for (x = 0; x < SIZE; x++)  
{  
    y += 5;  
    A[x] = calculation(y);  
}
```



```
int x;  
float A[SIZE];  
#pragma omp parallel  
{  
    #pragma omp for  
    for (x = 0; x < SIZE; x++)  
    {  
        int y = 10 + 5(x + 1);  
        A[x] = calculation(y);  
    }  
}
```

- ✓ Loop carried dependencies are no-go for parallelization.
- ✓ You have to rewrite the loop and eliminate the loop carried dependency.
- ✓ Rewrite the code in terms of the control variable.

OpenMP – Loop carried dependencies

Recognize this?

```
for (i = 0; i < steps; i++)  
{  
    x = (i + 0.5) * base;  
    fdx = 4 / (1 + x * x);  
    acum += fdx;  
}
```

← Loop carried dependency!

These loop carried dependencies have a name:

Reductions!



OpenMP has
built-in
REDUCTION!

Reduction

Syntax:
reduction (op:list)

A local copy of each variable in (list) is created and initialized depending on the operation (op).

At the end, all local copies are combined.



OpenMP – Worksharing – LOOP Construct - Reductions

Syntax:

reduction (op:list)

reduction (op:var1,var2)


Initial values for reduction operands.

Mathematically make sense.

Operator	Initialization
+	0
*	1
-	0
&	~0
	0
^	0
&&	1
	0

OpenMP – Worksharing – LOOP Construct - Reductions

```
for (i = 0; i < steps; i++)  
{  
    x = (i + 0.5) * base;  
    fdx = 4 / (1 + x * x);  
    acum += fdx;  
}
```



```
#pragma omp parallel  
{  
    #pragma omp for reduction (+:acum)  
    for (i = 0; i < steps; i++)  
    {  
        x = (i + 0.5) * base;  
        fdx = 4 / (1 + x * x);  
        acum += fdx;  
    }  
}
```



Exercise

1.- Instead of manually handling the threads, insert an OMP Loop Construct and let the compiler do the work sharing.

```
#include <stdio.h>
#include <time.h>

long cantidadIntervalos = 100000000;
double baseIntervalo;
double fdx;
double acum = 0;
clock_t start, end;

void main() {
    double x=0;
    long i;
    baseIntervalo = 1.0 / cantidadIntervalos;
    start = clock();
    for (i = 0; i < cantidadIntervalos; i++) {
        x = (i+0.5)*baseIntervalo;
        fdx = 4 / (1 + x * x);
        acum += fdx;
    }
    acum *= baseIntervalo;
    end = clock();
    printf("Result = %20.18lf (%ld)\n", acum, end - start);
}
```

- Use the original serial program.
- Parallelize it with an OMP LOOP Construct.
- You will use:
 - OpenMP PARALLEL directive.
 - #pragma omp parallel
 - FOR clause
 - #pragma omp for
 - REDUCTION on the accumulator.
 - reduction (op:list)
- Minimize the number of changes made.



Exercise – Code used

```
#include <stdio.h>
#include <time.h>
```

```
long cantidadIntervalos = 1000000000;
double baseIntervalo;
//double fdx;
double acum = 0;
```

Variable moved. Has to be private to each thread.

```
clock_t start, end;
```

```
void main() {
```

```
    //double x = 0; //Has to be local to the threads.
```

Variable moved

```
    long i;
    baseIntervalo = 1.0 / cantidadIntervalos;
    start = clock();
```



Exercise – Code used

```
#pragma omp parallel
```

Directive Inserted

```
{
```

```
double fdx; //Has to be local to the threads.
```

Variable Moved

```
double x; //Has to be local to the threads.
```

Variable Moved

```
#pragma omp for reduction(+:acum)
```

Directive Inserted

```
for (i = 0; i < cantidadIntervalos; i++) {
```

```
    x = (i + 0.5) * baseIntervalo;
```

```
    fdx = 4 / (1 + x * x);
```

```
    acum += fdx;
```

```
}
```

```
}
```

```
acum *= baseIntervalo;
```

```
end = clock();
```

```
printf("Result = %20.18lf (%1d)\n", acum, end - start);
```

```
}
```

OpenMP - Shared and Private Data **IMPLICIT**

A note on Shared and Private Data...

```
#include <stdio.h>
```

```
int main() {
```

```
    type global_variable = value;
```

```
    #pragma omp parallel
    { //Parallel region begins
```


```
        type private_variable = value;
```

```
        printf("Hello OMP World!...\n");
```

```
    } //Parallel region ends
```

```
    printf("This will print after sync...
    And only once!...\n");
```

```
}
```



OpenMP - Shared and Private Data **EXPLICIT**

```
#include <stdio.h>
```

```
int main() {
```

```
    int i, n, a, b;
```

```
    #pragma omp parallel shared (n, a) private (i, b)
    { //Parallel region begins
```

```
        for (i = 0; i < n; i++) {
            b = a + i;
        }
```

```
    } //Parallel region ends
```

```
    printf("This will print after sync...
    And only once!...\n");
```

```
}
```

Explicit
SHARED

Explicit
PRIVATE



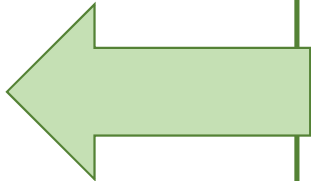
OpenMP – Example using PRIVATE

```
#include <stdio.h>
#include <time.h>

long cantidadIntervalos = 1000000000;
double baseIntervalo, fdx, acum = 0; //fdx will be local
clock_t start, end;

void main() {

    double x = 0; //Will be local, using PRIVATE.
    long i;
    baseIntervalo = 1.0 / cantidadIntervalos;
    start = clock();
```



OpenMP – Example using PRIVATE


```
#pragma omp parallel private (fdx,x)
{
    #pragma omp for reduction(+:acum)
    for (i = 0; i < cantidadIntervalos; i++) {
        x = (i + 0.5) * baseIntervalo;
        fdx = 4 / (1 + x * x);
        acum += fdx;
    }
    acum *= baseIntervalo;
    end = clock();
    printf("Result= %20.18lf (%ld)\n", acum, end-start);
}
```

fdx, x local to each thread.

OpenMP – #pragma omp parallel & #pragma omp for

You can merge both constructs into one:

```
#pragma omp parallel  
{  
    #pragma omp for  
    for (i = 0; i < 100000; i++)  
        A[i] += B[i];  
}
```



```
#pragma omp parallel for  
for (i = 0; i < 100000; i++)  
    A[i] += B[i];
```



This is true if `#pragma omp for` is the ONLY thing inside the brackets of `#pragma omp parallel`.



Exercise

1.- Use only ONE pragma.

```
#include <stdio.h>
#include <time.h>

long cantidadIntervalos = 1000000000;
double baseIntervalo;
double fdx;
double acum = 0;
clock_t start, end;

void main() {
    double x=0;
    long i;
    baseIntervalo = 1.0 / cantidadIntervalos;
    start = clock();
    for (i = 0; i < cantidadIntervalos; i++) {
        x = (i+0.5)*baseIntervalo;
        fdx = 4 / (1 + x * x);
        acum += fdx;
    }
    acum *= baseIntervalo;
    end = clock();
    printf("Result = %20.18lf (%ld)\n", acum, end - start);
}
```

- Use the original serial program.
- Parallelize it with an OMP LOOP Construct.
- You will use:
 - OpenMP PARALLEL directive.
 - FOR clause
 - REDUCTION on the accumulator.
- Minimize the number of changes made.



Exercise – Code used

```
#include <stdio.h>
#include <time.h>

long cantidadIntervalos = 1000000000;
double baseIntervalo;
double fdx;
double acum = 0;
clock_t start, end;


void main() {
    double x = 0;
    long i;
    baseIntervalo = 1.0 / cantidadIntervalos;
    start = clock();

    #pragma omp parallel for reduction (+:acum) private (fdx, x)
    for (i = 0; i < cantidadIntervalos; i++) {
        x = (i + 0.5) * baseIntervalo;
        fdx = 4 / (1 + x * x);
        acum += fdx;
    }
    acum *= baseIntervalo;
    end = clock();
    printf("Result = %20.18lf (%ld)\n", acum, end - start);
}
```

← OMP Directive + Clauses



OpenMP – Worksharing – LOOP CONSTRUCT - Schedule

- Compilers are “dumb”*.
- Programmer has to help.
- The problem... How to distribute the work the most efficient way possible?
- `#pragma omp parallel for` **`schedule(kind [,chunk size])`**

Optional
- If SCHEDULE clause is not present, the default Schedule is implementation-defined.

* Not really, they always play safe.



OpenMP – Worksharing – LOOP CONSTRUCT - Schedule

➤ #pragma omp parallel for **schedule(kind [,chunk size])**

Schedule **kind**:

Kind	Description
static	Divide the loop into equal-sized chunks or as equal as possible in the case where the number of loop iterations is not evenly divisible by the number of threads multiplied by the chunk size. By default, chunk size is loop_count/number_of_threads. Set chunk to 1 to interleave the iterations.
dynamic	Use the internal work queue to give a chunk-sized block of loop iterations to each thread. When a thread is finished, it retrieves the next block of loop iterations from the top of the work queue. By default, the chunk size is 1. Be careful when using this scheduling type because of the extra overhead involved.
guided	Similar to dynamic scheduling, but the chunk size starts off large and decreases to better handle load imbalance between iterations. The optional chunk parameter specifies them minimum size chunk to use. By default the chunk size is approximately loop_count/number_of_threads.



OpenMP – Worksharing – LOOP CONSTRUCT - Schedule

Static vs Dynamic... when to use each?

STATIC

- Pre-determined or known by the programmer.
- Least work at runtime.
- Scheduling done at **COMPILE time.**

DYNAMIC

- Unpredictable, highly variable work per iteration.
- Most work at runtime.
- Complex scheduling done at **RUNTIME.**



OpenMP – Worksharing – LOOP CONSTRUCT - Schedule

Another schedule **kind**: runtime

Kind	Description
static	Divide the loop into equal-sized chunks or as equal as possible in the case where the number of loop iterations is not evenly divisible by the number of threads multiplied by the chunk size. By default, chunk size is <code>loop_count/number_of_threads</code> . Set chunk to 1 to interleave the iterations.
dynamic	Use the internal work queue to give a chunk-sized block of loop iterations to each thread. When a thread is finished, it retrieves the next block of loop iterations from the top of the work queue. By default, the chunk size is 1. Be careful when using this scheduling type because of the extra overhead involved.
guided	Similar to dynamic scheduling, but the chunk size starts off large and decreases to better handle load imbalance between iterations. The optional chunk parameter specifies them minimum size chunk to use. By default the chunk size is approximately <code>loop_count/number_of_threads</code> .
runtime	Uses the <code>OMP_schedule</code> environment variable to specify which one of the three loop-scheduling types should be used. <code>OMP_SCHEDULE</code> is a string formatted exactly the same as would appear on the parallel construct.

No chunk_size

OpenMP – Worksharing – LOOP CONSTRUCT - Schedule

- Can change schedule without recompiling.
 - #pragma omp parallel for **schedule(runtime)**
 - Uses environment variable **OMP_SCHEDULE**
 - **set OMP_SCHEDULE=kind,chunk** ➤ **export OMP_SCHEDULE=kind,chunk**
 - **Ej: set OMP_SCHEDULE=static,1000** ➤ **Ej: export OMP_SCHEDULE=static,1000**
- | | | | |
|---|---|--|---|
|  | set OMP_SCHEDULE=dynamic
set OMP_SCHEDULE=guided,4 |  | export OMP_SCHEDULE=dynamic
export OMP_SCHEDULE=guided,4 |
|---|---|--|---|

Remarks

The default value in the Visual C++ implementation of the OpenMP standard is
`OMP_SCHEDULE=static,0.`

OpenMP – Revisiting Barriers – NOWAIT (**EXPLICIT**)


OpenMP Synchronization - **Barrier**

- No thread can proceed past a barrier **until all** the other threads have arrived.
- Syntax C/C++:

```
#pragma omp barrier
```

- Either all threads or none must encounter the barrier: otherwise

DEADLOCK!!



Allows us to place **EXPLICIT** barriers where needed.

OpenMP – Revisiting Barriers – NOWAIT (IMPLICIT)

```
#pragma omp parallel
{
    int threadID = omp_get_thread_num();
    A[threadID] = functionX(threadID);
    #pragma omp barrier
    #pragma omp for
    for (i = 0; i < SIZE; i++)
        C[threadID] = functionY(threadID, A);
    #pragma omp for nowait
    for (i = 0; i < SIZE; i++)
        B[threadID] = functionZ(threadID, C);
    A[threadID] = functionAA(threadID);
}
```

Explicit Barrier

Implicit Barrier

Implicit Barrier (overridden)

Implicit Barrier at end of parallel region (CANNOT be overridden)

Remember... Synchronization is EXPENSIVE!




If we can skip a synchronization, it is worth it!



If we as programmers are sure that our ALGORITHM allows us NOT to do a synchronization in a `#pragma omp for`:

`#pragma omp for nowait`



OpenMP – MASTER Construct

The **master** construct:

- Specifies a structured block that is executed by the master thread of the team.
- The syntax of the **master** directive is as follows:

```
#pragma omp master new-line  
structured-block
```

- Other threads in the team do not execute the associated structured block.
- There is no implied barrier either on entry to or exit from the master construct.



OpenMP – MASTER Construct

```
#include <omp.h>
#include <stdio.h>
#include <windows.h>

int main() {
    int sharedVar = 6;
    #pragma omp parallel num_threads(10)
    {
        int privateVar = omp_get_thread_num(); //This data is mine!
        sharedVar = omp_get_thread_num(); //I'll leave my mark!
        Sleep(2);

        #pragma omp master //Only 1 print of the final value by the Master Thread
        {
            printf("Thread (%d), privateVar = %d, sharedVar= %d \n",
                omp_get_thread_num(), privateVar, sharedVar);
        }
    }
}
```



OpenMP – SINGLE Construct

The **single** construct:

- Specifies that the associated structured block is executed by only one thread in the team.
- Not necessarily the master thread .
- The syntax of the **single** directive is as follows:

```
#pragma omp single [clause[,] clause] ...] new-line  
structured-block
```

- There is an implicit barrier after the **single** construct.
- Allows the use of `nowait`.



OpenMP – SINGLE Construct

```
#include <omp.h>
#include <stdio.h>
#include <windows.h>

int main() {
    int sharedVar = 6;
    #pragma omp parallel num_threads(10)
    {
        int privateVar = omp_get_thread_num(); //This data is mine!
        sharedVar = omp_get_thread_num(); //I'll leave my mark!
        Sleep(2);

        #pragma omp single //Any thread, print the final value
        {
            printf("Thread (%d), privateVar = %d, sharedVar= %d \n",
                omp_get_thread_num(), privateVar, sharedVar);
        }
    }
}
```




Exercise

1.- Play with the different scheduling options. What are your conclusions?

```
#include <stdio.h>
#include <time.h>

long cantidadIntervalos = 100000000;
double baseIntervalo;
double fdx;
double acum = 0;
clock_t start, end;

void main() {
    double x = 0;
    long i;
    baseIntervalo = 1.0 / cantidadIntervalos;
    start = clock();

    #pragma omp parallel for reduction (+:acum) private (fdx, x)
    for (i = 0; i < cantidadIntervalos; i++) {
        x = (i + 0.5) * baseIntervalo;
        fdx = 4 / (1 + x * x);
        acum += fdx;
    }
    acum *= baseIntervalo;
    end = clock();
    printf("Result = %20.18lf (%ld)\n", acum, end - start);
}
```

- Use the OMP code with just 1 directive.
 - Use the SCHEDULE clause with:
 - schedule (static,chunk size)
 - schedule (dynamic,chunk size)
 - schedule (guided,chunk size)
 - schedule (runtime)
 - set OMP_SCHEDULE=kind,chunk



OpenMP – Worksharing – LOOP CONSTRUCT - Schedule

```
#pragma omp parallel for reduction (+:acum) private (fdx, x) schedule (static)

#pragma omp parallel for reduction (+:acum) private (fdx, x) schedule (static, cantidadIntervalos)

#pragma omp parallel for reduction (+:acum) private (fdx, x) schedule (static, cantidadIntervalos/2)

#pragma omp parallel for reduction (+:acum) private (fdx, x) schedule (static, cantidadIntervalos/4)

#pragma omp parallel for reduction (+:acum) private (fdx, x) schedule (static, cantidadIntervalos/8)

#pragma omp parallel for reduction (+:acum) private (fdx, x) schedule (dynamic)

#pragma omp parallel for reduction (+:acum) private (fdx, x) schedule (dynamic, 1)

#pragma omp parallel for reduction (+:acum) private (fdx, x) schedule (dynamic, 2000)

#pragma omp parallel for reduction (+:acum) private (fdx, x) schedule (runtime) |set OMP_SCHEDULE=static,2

#pragma omp parallel for reduction (+:acum) private (fdx, x) schedule (runtime) | set OMP_SCHEDULE=guided,2
```



OpenMP – SECTIONS/SECTION Directive

The **sections/section** directive:


- Identifies a **noniterative** work-sharing construct that specifies a set of constructs that are to be divided among threads in a team.
- Each section is executed once by a thread in the team.

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        function_x();
        #pragma omp section
        function_y();
    }
}
```

- **There is an implicit barrier** after the **sections** construct.
- Allows the use of **nowait**.

OpenMP – SECTIONS/SECTION Construct Example

```
#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        function_x();
        #pragma omp section
        function_y();
        #pragma omp section
        function_z();
    }
}
```



OpenMP – #pragma omp parallel & #pragma omp sections

You can merge both constructs into one:

```
#pragma omp parallel  
{  
    #pragma omp sections  
    {  
        //Add  
        #pragma omp section  
    }  
}
```



```
#pragma omp parallel sections  
{  
    //Add  
    #pragma omp section  
}
```



This is true if `#pragma omp sections` is the ONLY thing inside the brackets of `#pragma omp parallel`.



For NEXT class...

Create a program using OMP to do task decomposition:

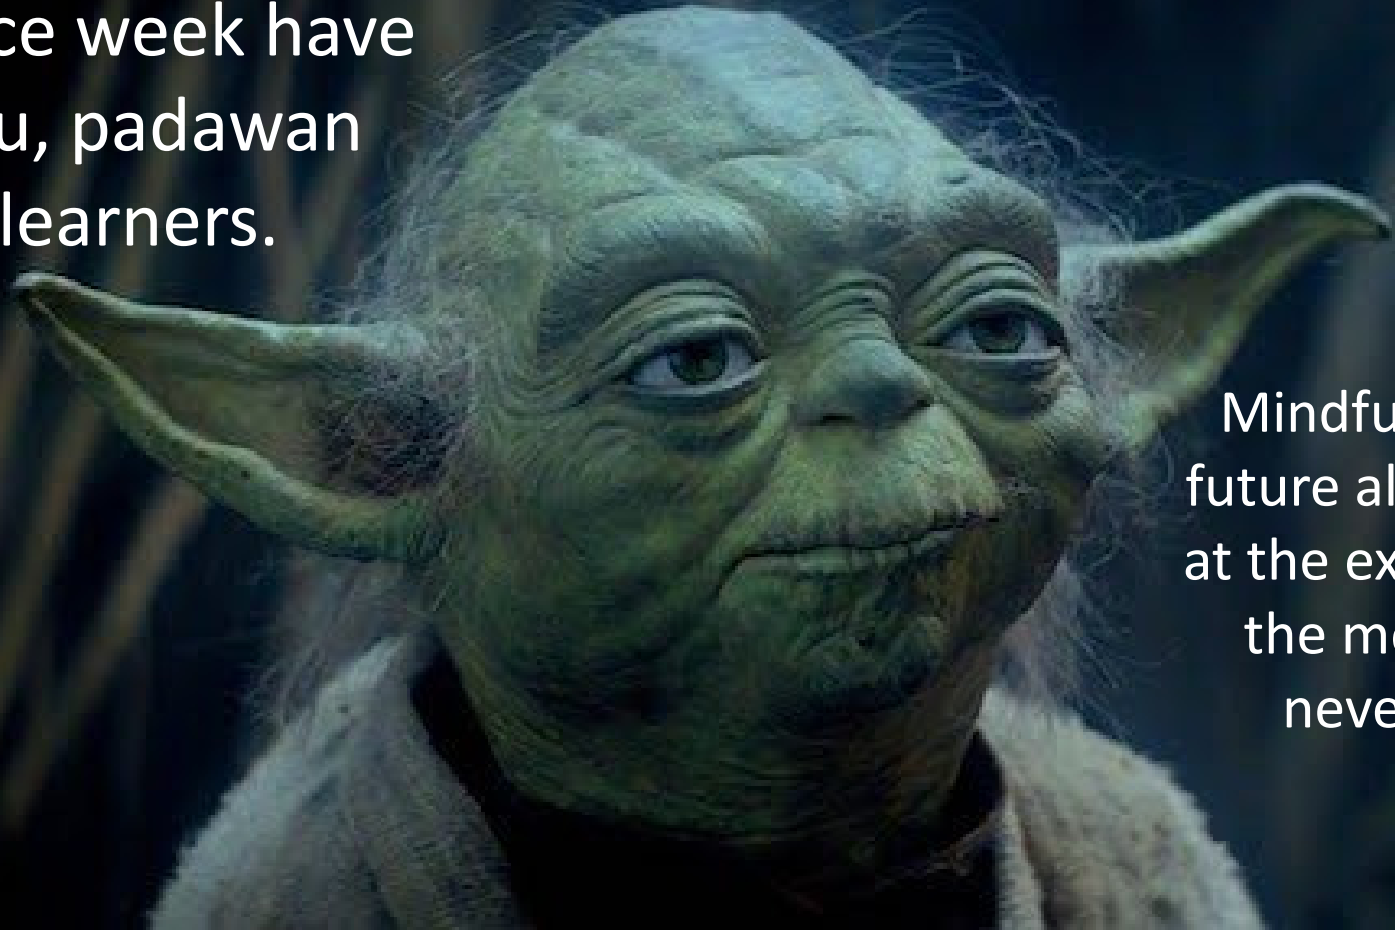
1.- Create a program that do the following:

- **Define 6 Arrays of FLOATS.**
 - **16 elements each.**
 - **Array A = {10.0, 11.0, 12.0... 25.0}**
 - **Array B = {1.0, 2.0, 3.0... 16.0}**
 - **Array C, D, E & F are zeroed at the beginning.**
 - **Using OMP, do task decomposition:**
 - **$C[i] = A[i] + B[i];$**
 - **$D[i] = A[i] - B[i];$**
 - **$E[i] = A[i] * B[i];$**
 - **$F[i] = A[i] / B[i];$**
 - **Print the 4 lines with the results. One line per result.**
- 2.- Hand in your code and a screenshot of your output.



See you soon...

A nice week have
you, padawan
learners.



Mindful of the
future always be,
at the expense of
the moment
never be.