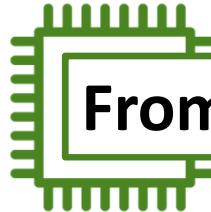


Feb-Jun 2021

# Multiprocessors

## Open Multi-Processing

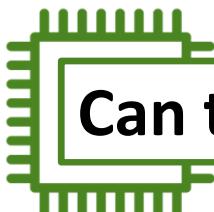


## From Homework...

Create a program using OMP to do task decomposition:

1.- Create a program that do the following:

- **Define 6 Arrays of FLOATS.**
- **16 elements each.**
  - **Array A = {10.0, 11.0, 12.0... 25.0}**
  - **Array B= {1.0, 2.0, 3.0... 16.0}**
  - **Array C, D, E & F are zeroed at the begining.**
- **Using OMP, do task decomposition:**
  - **C[i] = A[i] + B[i];**
  - **D[i] = A[i] - B[i];**
  - **E[i] = A[i] \* B[i];**
  - **F[i] = A[i] / B[i];**
- **Print the 4 lines with the results. One line per result.**
- **Hand in your code and a screenshot of your output.**



## Can this be improved (1)?

```
#pragma omp parallel num_threads(4)
{
    for (int i = 0; i < 16; i++)
    {
        if (omp_get_thread_num() == 0)
            C[i] = A[i] + B[i];

        if (omp_get_thread_num() == 1)
            D[i] = A[i] - B[i];

        if (omp_get_thread_num() == 2)
            E[i] = A[i] * B[i];

        if (omp_get_thread_num() == 3)
            F[i] = A[i] / B[i];
    }
}
```

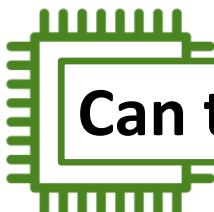


```
#pragma omp parallel num_threads(4)
{
    if (omp_get_thread_num() == 0)
        for (int i = 0; i < 16; i++)
            C[i] = A[i] + B[i];

    if (omp_get_thread_num() == 1)
        for (int i = 0; i < 16; i++)
            D[i] = A[i] - B[i];

    if (omp_get_thread_num() == 2)
        for (int i = 0; i < 16; i++)
            E[i] = A[i] * B[i];

    if (omp_get_thread_num() == 3)
        for (int i = 0; i < 16; i++)
            F[i] = A[i] / B[i];
}
```



## Can this be improved (2)?

```
#pragma omp parallel num_threads(4)
{
    if (omp_get_thread_num() == 0)
        for (int i = 0; i < 16; i++)
            C[i] = A[i] + B[i];

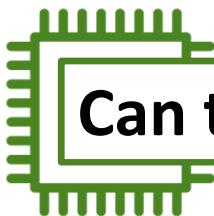
    if (omp_get_thread_num() == 1)
        for (int i = 0; i < 16; i++)
            D[i] = A[i] - B[i];

    if (omp_get_thread_num() == 2)
        for (int i = 0; i < 16; i++)
            E[i] = A[i] * B[i];

    if (omp_get_thread_num() == 3)
        for (int i = 0; i < 16; i++)
            F[i] = A[i] / B[i];
}
```



```
#pragma omp parallel
{
    int thread_num = omp_get_thread_num();
    switch (thread_num)
    {
        case 0: //Add
            for (int i = 0; i < 16; i++) {
                arrayC[i] = arrayA[i] + arrayB[i];
            };
            break;
        case 1: //Subtract
            for (int i = 0; i < 16; i++) {
                arrayD[i] = arrayA[i] - arrayB[i];
            };
            break;
        case 2: //Multiply
            for (int i = 0; i < 16; i++) {
                arrayE[i] = arrayA[i] * arrayB[i];
            };
            break;
        case 3: //Divide
            for (int i = 0; i < 16; i++) {
                arrayF[i] = arrayA[i] / arrayB[i];
            };
            break;
    }
}
```

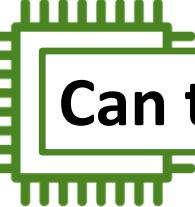


## Can this be improved (3)?

```
#pragma omp parallel
{
    int thread_num = omp_get_thread_num();
    switch (thread_num)
    {
        case 0: //Add
            for (int i = 0; i < 16; i++) {
                arrayC[i] = arrayA[i] + arrayB[i];
            }
            break;
        case 1: //Subtract
            for (int i = 0; i < 16; i++) {
                arrayD[i] = arrayA[i] - arrayB[i];
            }
            break;
        case 2: //Multiply
            ...
        case 3: //Divide
            ...
    }
}
```



```
#pragma omp parallel
{
    int thread_num = omp_get_thread_num();
    switch (thread_num)
    {
        case 0: //Add
            for (int i = 0; i < 16/4; i++) {
                __m128 a = _mm128_load_ps(&A[i*4]);
                __m128 b = _mm128_load_ps(&B[i*4]);
                __m128 sum = _mm128_add_ps(a, b);
                _mm128_store_ps(&C[i*4], sum);
            }
            break;
        case 1: //Subtract
            for (int i = 0; i < 16/4; i++) {
                __m128 a = _mm128_load_ps(&A[i*4]);
                __m128 b = _mm128_load_ps(&B[i*4]);
                __m128 sum = _mm128_sub_ps(a, b);
                _mm128_store_ps(&C[i*4], sum);
            }
            break;
        case 2: //Multiply
            ...
        case 3: //Divide
            ...
    }
}
```

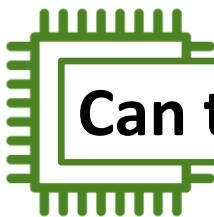


## Can this be improved (4)?

```
#pragma omp parallel
{
    int thread_num = omp_get_thread_num();
    switch (thread_num)
    {
        case 0: //Add
            for (int i = 0; i < 16/4; i++) {
                __m128 a = _mm128_load_ps(&A[i*4]);
                __m128 b = _mm128_load_ps(&B[i*4]);
                __m128 sum = _mm128_add_ps(a, b);
                _mm128_store_ps(&C[i*4], sum);;
            };
            break;
        case 1: //Subtract
            for (int i = 0; i < 16/4; i++) {
                __m128 a = _mm128_load_ps(&A[i*4]);
                __m128 b = _mm128_load_ps(&B[i*4]);
                __m128 sum = _mm128_sub_ps(a, b);
                _mm128_store_ps(&C[i*4], sum);;
            };
            break;
        case 2: //Multiply
            ...
        case 3: //Divide
            ...
    }
}
```



```
#pragma omp parallel
{
    int thread_num = omp_get_thread_num();
    switch (thread_num)
    {
        case 0: //Add
            for (int i = 0; i < 16/8; i++) {
                __m256 a = _mm256_load_ps(&A[i * 8]);
                __m256 b = _mm256_load_ps(&B[i * 8]);
                __m256 sum = _mm256_add_ps(a, b);
                _mm256_store_ps(&C[i * 8], sum);
            };
            break;
        case 1: //Subtract
            for (int i = 0; i < 16/8; i++) {
                __m256 a = _mm256_load_ps(&A[i * 8]);
                __m256 b = _mm256_load_ps(&B[i * 8]);
                __m256 sum = _mm256_sub_ps(a, b);
                _mm256_store_ps(&C[i * 8], sum);
            };
            break;
        case 2: //Multiply
            ...
        case 3: //Divide
            ...
    }
}
```

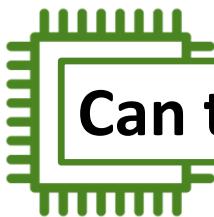


## Can this be improved (5)?

```
#pragma omp parallel
{
    int thread_num = omp_get_thread_num();
    switch (thread_num)
    {
        case 0: //Add
            for (int i = 0; i < 16/8; i++) {
                __m256 a = _mm256_load_ps(&A[i * 8]);
                __m256 b = _mm256_load_ps(&B[i * 8]);
                __m256 sum = _mm256_add_ps(a, b);
                _mm256_store_ps(&C[i * 8], sum);
            };
            break;
        case 1: //Subtract
            for (int i = 0; i < 16/8; i++) {
                __m256 a = _mm256_load_ps(&A[i * 8]);
                __m256 b = _mm256_load_ps(&B[i * 8]);
                __m256 sum = _mm256_sub_ps(a, b);
                _mm256_store_ps(&C[i * 8], sum);
            };
            break;
        case 2: //Multiply
            ...
        case 3: //Divide
            ...
    }
}
```



```
#pragma omp parallel
{
    int thread_num = omp_get_thread_num();
    switch (thread_num)
    {
        case 0: //Add
            __m256 a = _mm256_load_ps(&A[0]);
            __m256 b = _mm256_load_ps(&B[0]);
            __m256 sum = _mm256_add_ps(a, b);
            _mm256_store_ps(&C[0], sum);
            a = _mm256_load_ps(&A[8]);
            b = _mm256_load_ps(&B[8]);
            sum = _mm256_add_ps(a, b);
            _mm256_store_ps(&C[8], sum);
        };
        break;
        case 1: //Subtract
            ...
        case 2: //Multiply
            ...
        case 3: //Divide
            ...
    }
}
```

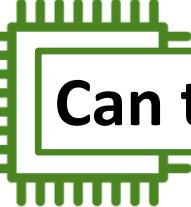


## Can this be improved (6)?

```
#pragma omp parallel
{
    int thread_num = omp_get_thread_num();
    switch (thread_num)
    {
        case 0: //Add
        {
            __m256 a = _mm256_load_ps(&A[0]);
            __m256 b = _mm256_load_ps(&B[0]);
            __m256 sum = _mm256_add_ps(a, b);
            _mm256_store_ps(&C[0], sum);
            a = _mm256_load_ps(&A[8]);
            b = _mm256_load_ps(&B[8]);
            sum = _mm256_add_ps(a, b);
            _mm256_store_ps(&C[8], sum);
        };
        break;
        case 1: //Subtract
        ...
        case 2: //Multiply
        ...
        case 3: //Divide
        ...
    }
}
```

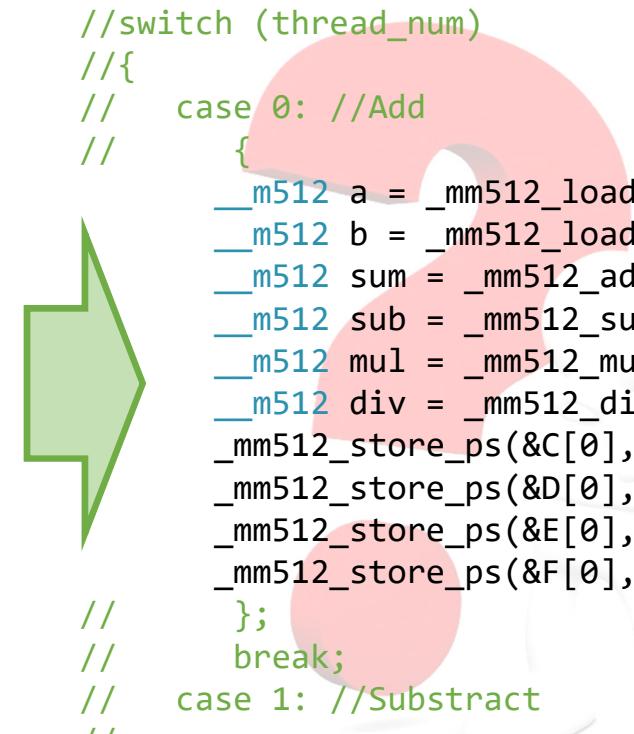


```
#pragma omp parallel
{
    int thread_num = omp_get_thread_num();
    switch (thread_num)
    {
        case 0: //Add
        {
            __m512 a = _mm512_load_ps(&A[0]);
            __m512 b = _mm512_load_ps(&B[0]);
            __m512 sum = _mm512_add_ps(a, b);
            _mm512_store_ps(&C[0], sum);
        };
        break;
        case 1: //Subtract
        ...
        case 2: //Multiply
        ...
        case 3: //Divide
        ...
    }
}
```

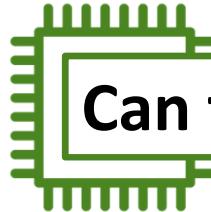


## Can this be improved (7)?

```
#pragma omp parallel
{
    int thread_num = omp_get_thread_num();
    switch (thread_num)
    {
        case 0: //Add
        {
            __m512 a = _mm512_load_ps(&A[0]);
            __m512 b = _mm512_load_ps(&B[0]);
            __m512 sum = _mm512_add_ps(a, b);
            _mm512_store_ps(&C[0], sum);
        };
        break;
        case 1: //Substract
        ...
        case 2: //Multiply
        ...
        case 3: //Divide
        ...
    }
}
```



```
//#pragma omp parallel
//{
//int thread_num = omp_get_thread_num();
//switch (thread_num)
//{
//    case 0: //Add
//    {
//        __m512 a = _mm512_load_ps(&A[0]);
//        __m512 b = _mm512_load_ps(&B[0]);
//        __m512 sum = _mm512_add_ps(a, b);
//        __m512 sub = _mm512_sub_ps(a, b);
//        __m512 mul = _mm512_mul_ps(a, b);
//        __m512 div = _mm512_div_ps(a, b);
//        _mm512_store_ps(&C[0], sum);
//        _mm512_store_ps(&D[0], sub);
//        _mm512_store_ps(&E[0], mul);
//        _mm512_store_ps(&F[0], div);
//    };
//    break;
//    case 1: //Substract
//    ...
//    case 2: //Multiply
//    ...
//    case 3: //Divide
//    ...
//}
```

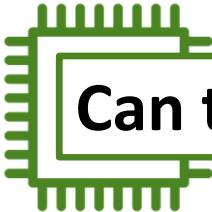


## Can this be improved (7)?

- In this case, consider thread creation costs vs. 6 more intrinsics executed by the main thread...

```
_m512 a = _mm512_load_ps(&A[0]);
_m512 b = _mm512_load_ps(&B[0]);
_m512 sum = _mm512_add_ps(a, b);
_m512 sub = _mm512_sub_ps(a, b); //1
_m512 mul = _mm512_mul_ps(a, b); //2
_m512 div = _mm512_div_ps(a, b); //3
_mm512_store_ps(&C[0], sum);
_mm512_store_ps(&D[0], sub);      //4
_mm512_store_ps(&E[0], mul);     //5
_mm512_store_ps(&F[0], div);     //6
```

- For low repeat cycles, Intrinsics will be faster.

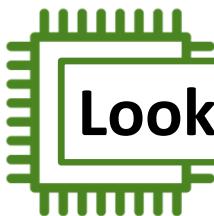


## Can this be improved (8)?

```
_m512 a = _mm512_load_ps(&A[0]);
_m512 b = _mm512_load_ps(&B[0]);
_m512 sum = _mm512_add_ps(a, b);
_m512 sub = _mm512_sub_ps(a, b);
_m512 mul = _mm512_mul_ps(a, b);
_m512 div = _mm512_div_ps(a, b);
_mm512_store_ps(&C[0], sum);
_mm512_store_ps(&D[0], sub);
_mm512_store_ps(&E[0], mul);
_mm512_store_ps(&F[0], div);
```



```
_m512 a = _mm512_load_ps(&A[0]);
_m512 b = _mm512_load_ps(&B[0]);
_mm512_store_ps(&C[0], _mm512_add_ps(a, b));
_mm512_store_ps(&D[0], _mm512_sub_ps(a, b));
_mm512_store_ps(&E[0], _mm512_mul_ps(a, b));
_mm512_store_ps(&F[0], _mm512_div_ps(a, b));
```



## Look at the Assembly...

```
case 0:  
{  
    __m256 a = _mm256_load_ps(&A[0]);  
    __m256 b = _mm256_load_ps(&B[0]);  
    __m256 sum = _mm256_add_ps(a, b);  
    _mm256_store_ps(&C[0], sum);  
}  
break;
```

```
case 0:  
{  
    __m256 a = _mm256_load_ps(&A[0]);  
    __m256 b = _mm256_load_ps(&B[0]);  
    _mm256_store_ps(&C[0], _mm256_add_ps(a, b));  
}  
break;
```

No more benefit.

```
case 0:  
{  
    __m256 a = _mm256_load_ps(&A[0]);  
    __m256 b = _mm256_load_ps(&B[0]);  
    eax,dword ptr [B]  
    esi  
    __m256 sum = _mm256_add_ps(a, b);  
    ymm0,ymm0,ymmmword ptr [eax]  
    vaddps  
    ymm0,ymm0,ymmmword ptr [eax]  
    _mm256_store_ps(&C[0], sum);  
    mov  
    eax,dword ptr [C]  
    vmovups  
    ymmword ptr [eax],ymm0  
    vzeroupper  
}  
break;  
  
00B11DEA mov  
00B11DED pop  
00B11DEE vmovups  
00B11DF2 mov  
00B11DF5 vaddps  
00B11DF9 mov  
00B11DFC vmovups  
00B11E00 vzeroupper  
  
case 0:  
{  
    __m256 a = _mm256_load_ps(&A[0]);  
    __m256 b = _mm256_load_ps(&B[0]);  
    eax,dword ptr [B]  
    esi  
    _mm256_store_ps(&C[0], _mm256_add_ps(a, b));  
    ymm0,ymmmword ptr [eax]  
    vaddps  
    ymm0,ymm0,ymmmword ptr [eax]  
    mov  
    eax,dword ptr [C]  
    vmovups  
    ymmword ptr [eax],ymm0  
    vzeroupper  
}  
break;
```



## Solution – Task Decomposition using OMP Sections

### 1.- Using OMP SECTIONS/SECTION....

```
#include <omp.h>
#include <stdio.h>
#include <intrin.h>

int main()
{
float __declspec(align(16)) arrayA[16] = { 10.0,11.0,12.0,13.0,14.0,15.0,16.0,17.0,18.0,19.0,20.0,21.0,22.0,23.0,24.0,25.0 };
float __declspec(align(16)) arrayB[16] = { 1.0,2.0,3.0,4.0,5.0,6.0,7.0,8.0,9.0,10.0,11.0,12.0,13.0,14.0,15.0,16.0 };
float __declspec(align(16)) arrayC[16], arrayD[16], arrayE[16], arrayF[16] = { 0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0,0.0 };

#pragma omp parallel
{
    #pragma omp single
    {
        printf("Total threads in parallel region: %d\n\n", omp_get_num_threads());
    }

    #pragma omp sections
    {
        //Add
        #pragma omp section
        {
            printf("Thread in charge of addition: %d\n", omp_get_thread_num());
            for (int i = 0; i < 16; i++)
            {
                arrayC[i] = arrayA[i] + arrayB[i];
            }
        }

        //Subtract
        #pragma omp section
        {
            printf("Thread in charge of subtraction: %d\n", omp_get_thread_num());
            for (int i = 0; i < 16; i++)
            {
                arrayD[i] = arrayA[i] - arrayB[i];
            }
        }

        //Multiplication
        #pragma omp section
        {
            printf("Thread in charge of multiplication: %d\n", omp_get_thread_num());
            for (int i = 0; i < 16; i++)
            {
                arrayE[i] = arrayA[i] * arrayB[i];
            }
        }

        //Division
        #pragma omp section
        {
            printf("Thread in charge of division: %d\n", omp_get_thread_num());
            for (int i = 0; i < 16; i++)
            {
                arrayF[i] = arrayA[i] / arrayB[i];
            }
        }
    }
}
```

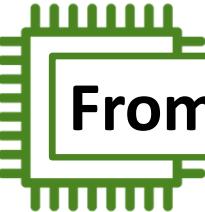


## Solution – Task Decomposition using OMP Sections

```
#pragma omp sections
{
    //Add
    #pragma omp section
    {
        printf("Thread in charge of addition: %d\n", omp_get_thread_num());
        for (int i = 0; i < 16; i++)
            arrayC[i] = arrayA[i] + arrayB[i];
    }

    //Subtract
    #pragma omp section
    {
        printf("Thread in charge of subtraction: %d\n", omp_get_thread_num());
        for (int i = 0; i < 16; i++)
            arrayD[i] = arrayA[i] - arrayB[i];
    }

    printf("Thread in charge of division: %d\n", omp_get_thread_num());
    for (int i = 0; i < 16; i++)
        arrayF[i] = arrayA[i] / arrayB[i];
}
}
```



## From Homework... Any question?

Create a program using OMP to do task decomposition:

1.- Create a program that do the following:

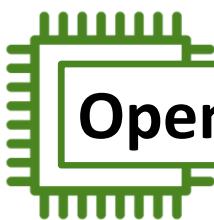
- **Define 6 Arrays of FLOATS.**
  - **16 elements each.**
    - **Array A = {10.0, 11.0, 12.0... 25.0}**
    - **Array B= {1.0, 2.0, 3.0... 16.0}**
    - **Array C, D, E & F are zeroed at the begining.**
  - **Using OMP, do task decomposition:**
    - **C[i] = A[i] + B[i];**
    - **D[i] = A[i] - B[i];**
    - **E[i] = A[i] \* B[i];**
    - **F[i] = A[i] / B[i];**
  - **Print the 4 lines with the results. One line per result.**
- 2.- Hand in your code and a screenshot of your output.



## OpenMP – BIG Summary

- So, by now, you have a pretty good idea of Open MP.
- We'll put some formallity now.





## OpenMP MACRO: `_OPENMP`

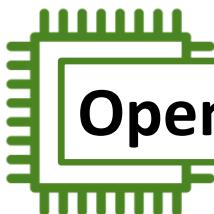
- Is defined by OpenMP-compliant implementations as the decimal constant `yyyymm`, which will be the year and month of the approved specification.
- How to know which OpenMP specification is implemented in the compiler?
- Just print the constant....

```
#include <omp.h>
#include <stdio.h>

int main()
{
    printf("OpenMP Version %d\n", _OPENMP);
    return 0;
}
```

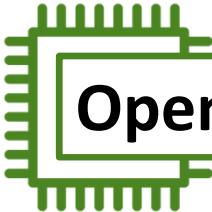
```
mp@multiprocesadores:~/multiprocs$ ./version
OpenMP Version 201511
mp@multiprocesadores:~/multiprocs$
```

Microsoft Visual Studio Debug Console  
OpenMP Version 200203  
C:\Users\Alex\source\repos\Te...  
To automatically close the con...  
le when debugging stops.  
Press any key to close this w...



## OpenMP Specifications for C/C++

	OpenMP Specification
199810	1.0
200203	2.0
200505	2.5
200811	3.0
201109	3.1
201310	4.0
201511	4.5
201811	5.0



## OpenMP Conditional Compilation

```
#include <omp.h>
#include <stdio.h>

int main()
{
    #ifdef _OPENMP
        int numThreads = omp_get_max_threads();
        printf("Number of threads %d\n\n", numThreads);
        omp_set_num_threads(numThreads - 4);
    #endif

    #pragma omp parallel
    { // Start of parallel region
        printf("Hello OMP World!...\\n");
    } // End of parallel region

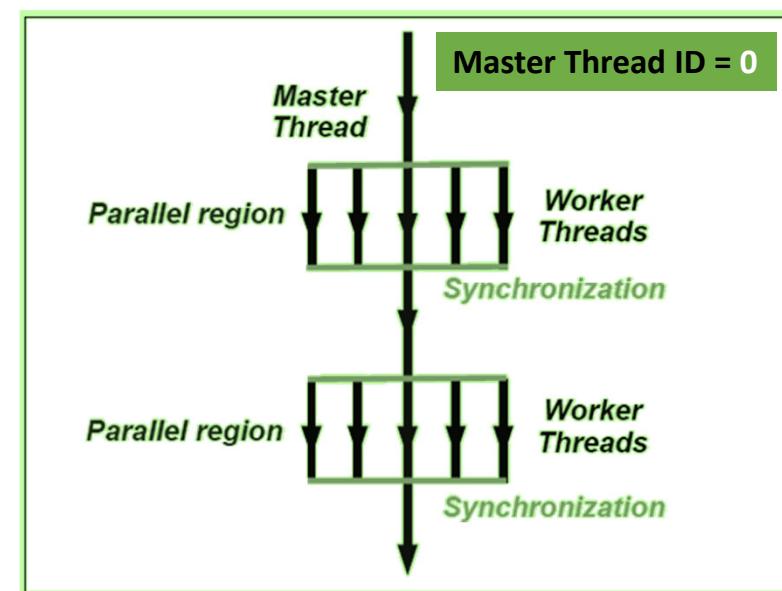
    return 0;
}
```

## Fork-join model

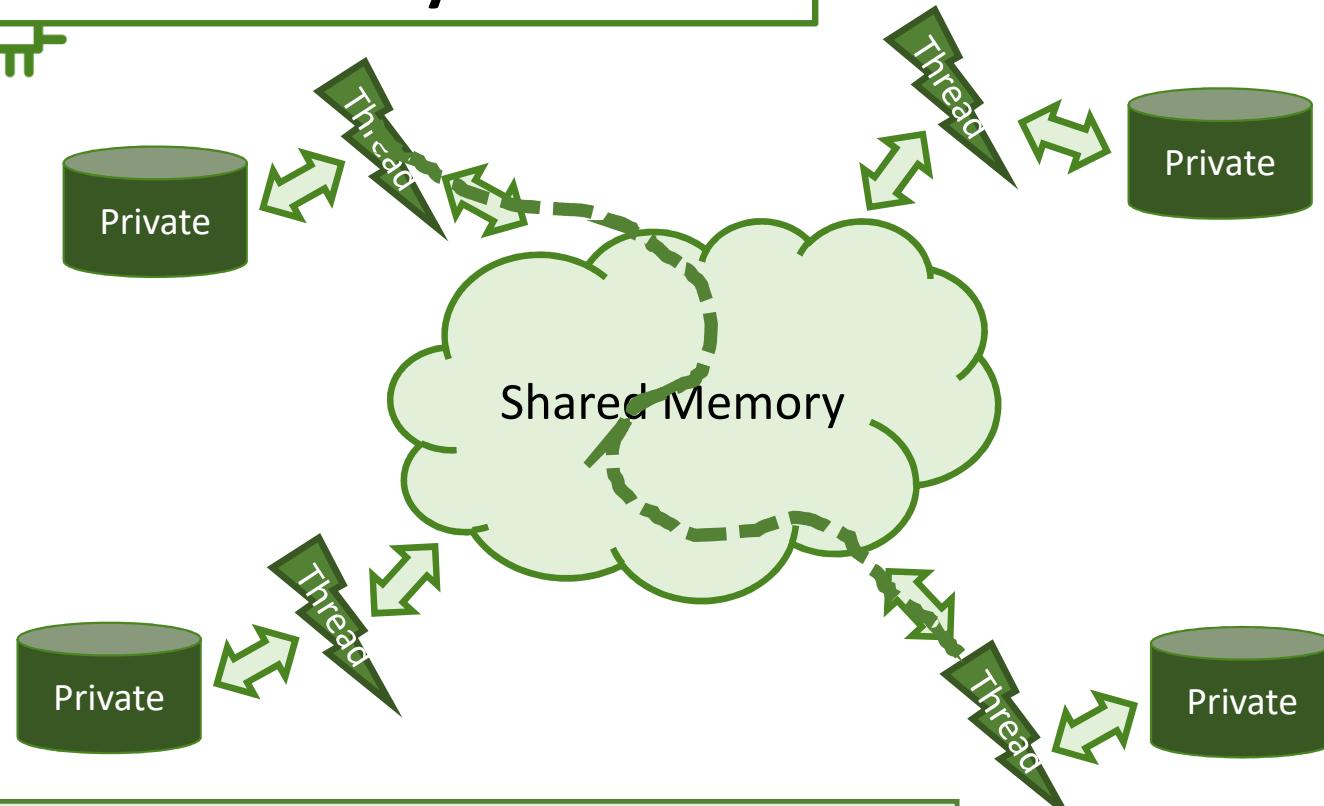
- The model uses one main thread and forks to multithread execution.
- The model may use several blocks of parallelization.

Fork = Parallel Region

Join = Synchronization



# OpenMP Memory Model

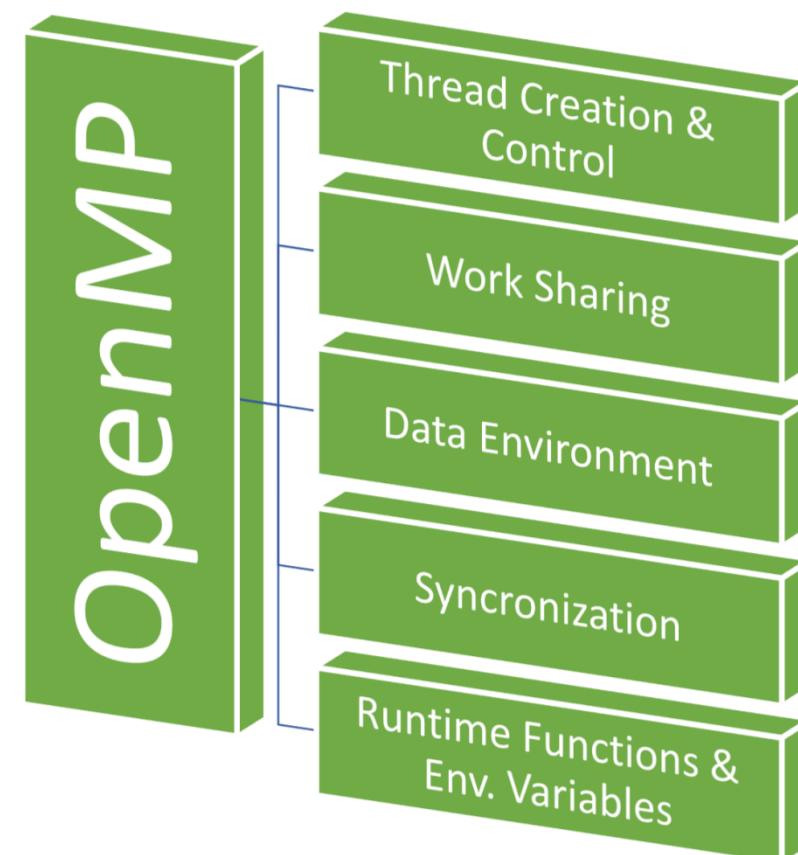


2 types of  
memory

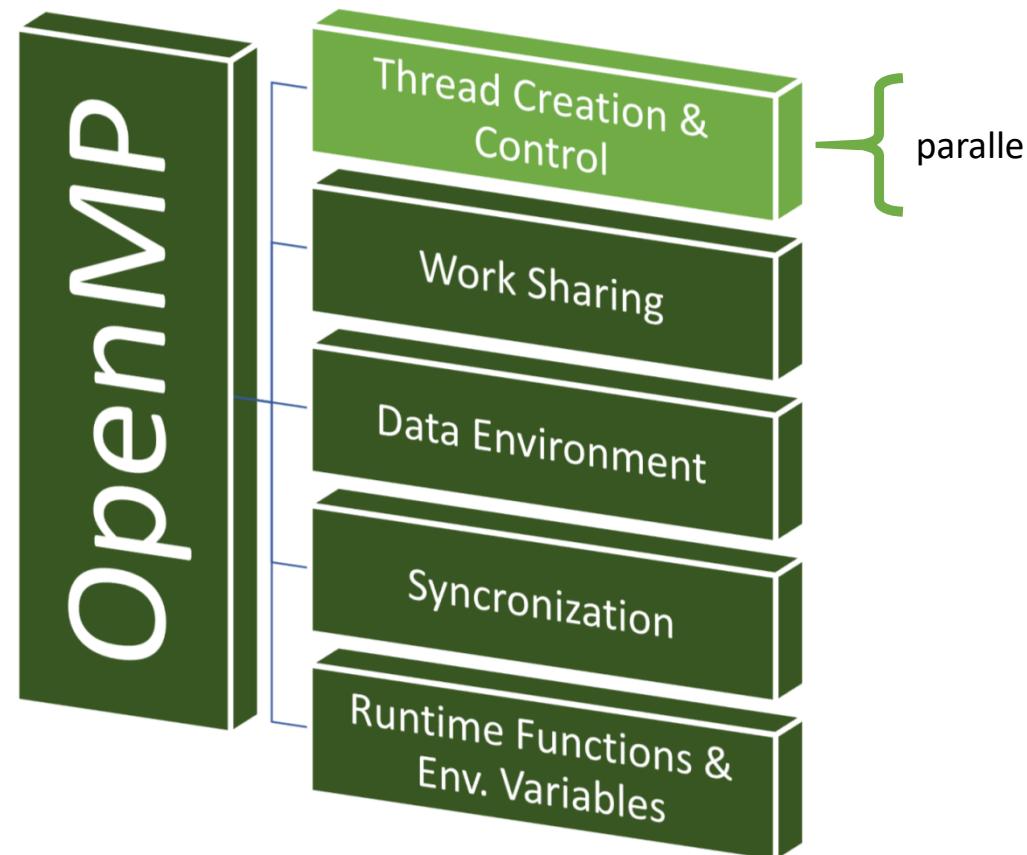
Heap – For all the program (OMP = Shared)

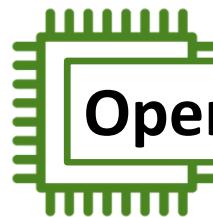
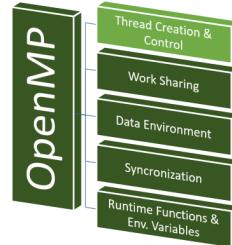
Stack – For just the thread (OMP = Private)

# OpenMP Core Elements



## OpenMP Core Elements





# OpenMP – Thread Creation & Control

## *parallel* construct

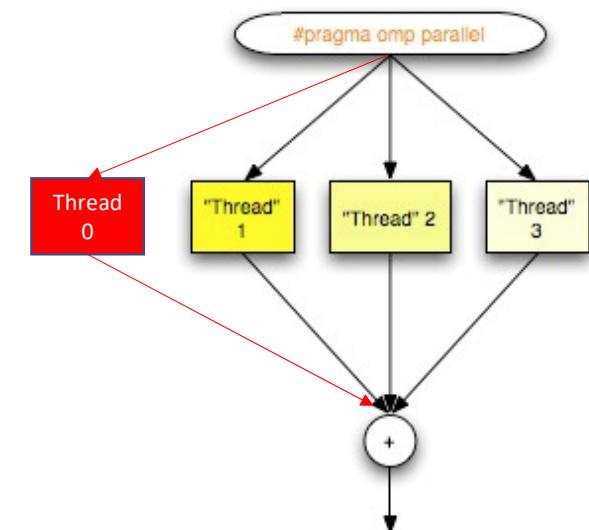
- Defines a parallel region, which is a region of the program that is to be executed by multiple threads in parallel.
- Is the fundamental construct that starts parallel execution.

```
#pragma omp parallel [clause[ [, ]clause] ...] new-line structured-block
```

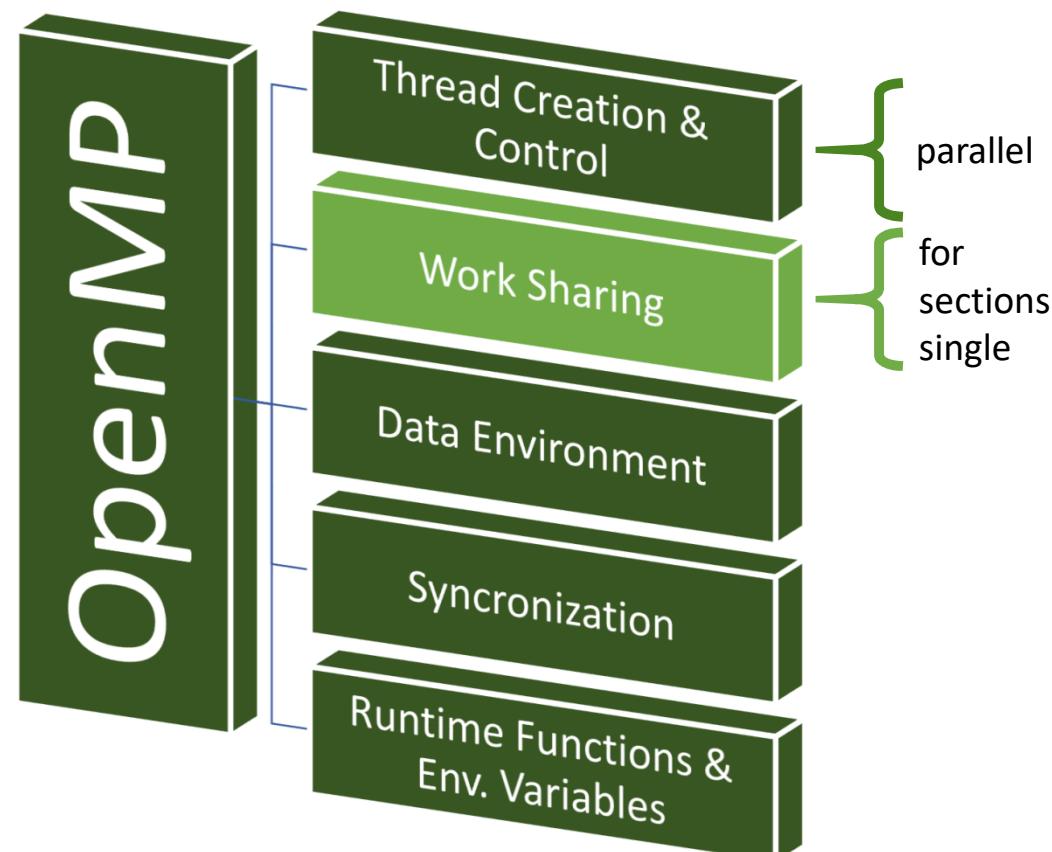
### Clauses

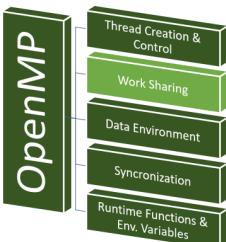
if(scalar-expression)  
 private(variable-list)  
 firstprivate(variable-list)  
 default(shared | none)  
 shared(variable-list)  
 copyin(variable-list)  
 reduction(operator: variable-list)  
 num\_threads(integer-expression)

```
#pragma omp parallel
{
  //Parallelized statements
}
```



# OpenMP Core Elements





# OpenMP – Work Sharing

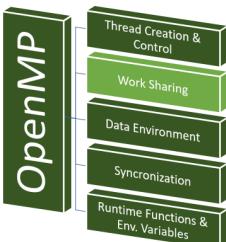
## for construct

- Identifies an iterative work-sharing construct that specifies that the iterations of the associated loop will be executed in parallel.
- The iterations of the **for** loop are distributed across threads that already exist in the team executing the parallel construct to which it binds.

```
#pragma omp for [clause[, clause] ... ] new-line for-loop
```

**Clauses**

- private(variable-list)
- firstprivate(variable-list)
- lastprivate(variable-list)
- reduction(operator: variable-list)
- ordered
- schedule(kind[, chunk\_size])
- nowait



# OpenMP – Work Sharing

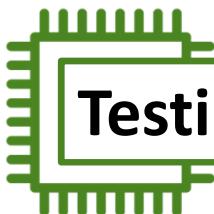
## ***sections*** construct

- The ***sections*** directive identifies a noniterative work-sharing construct that specifies a set of constructs that are to be divided among threads in a team. Each section is executed once by a thread in the team.

```
#pragma omp sections [clause[, clause] ...] new-line
{
    [#pragma omp section new-line]
    structured-block
    [#pragma omp section new-line
    structured-block ]
    ...
}
```

## Clauses

**private(variable-list)**  
**firstprivate(variable-list)**  
**lastprivate(variable-list)**  
**reduction(operator: variable-list)**  
**nowait**



## Testing OpenMP...

How many threads will execute a “section” section? Which ones?

**Write a small program to test it and answer the following questions:**

**1.- If my program has 4 sections and only 2 threads, which sections will be executed? Explain your answer.**

**2.- If my program has now 2 sections and 4 threads, what will happen?**

**3.- Can ThreadID=0 be assigned with a section to process or only ThreadID $\geq$ 1?**



## Testing OpenMP... Testing Code

```
#include <omp.h>
#include <stdio.h>
#include <Windows.h>

void main() {

    #pragma omp parallel num_threads(8)
    {
        #pragma omp single
        printf("Total threads available (%d): %d \n\n", omp_get_thread_num(), omp_get_num_threads());
        #pragma omp sections nowait
        {
            #pragma omp section
            {
                Sleep(10);
                printf("Thread (%d), Hello from section 1\n\n", omp_get_thread_num());
            }
            #pragma omp section
            {
                Sleep(1000);
                printf("Thread (%d), Hello from section 2\n\n", omp_get_thread_num());
            }
            #pragma omp section
            {
                Sleep(1000);
                printf("Thread (%d), Hello from section 3\n\n", omp_get_thread_num());
            }
            #pragma omp section
            {
                Sleep(1000);
                printf("Thread (%d), Hello from section 4\n\n", omp_get_thread_num());
            }
        }
        printf("Thread (%d) outside section\n", omp_get_thread_num());
    }
}
```

# OpenMP – Work Sharing

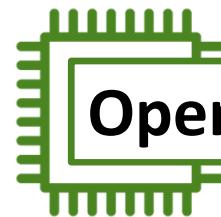
## **single** construct

- The **single** directive identifies a construct that specifies that the associated structured block is executed by only one thread in the team (not necessarily the master thread).

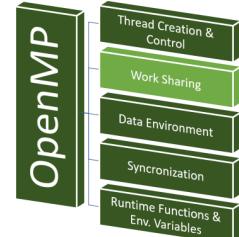
```
#pragma omp single [clause[, clause] ...] new-line
structured-block
```

### Clauses

```
private (variable-list)
firstprivate (variable-list)
copyprivate (variable-list)
nowait
```

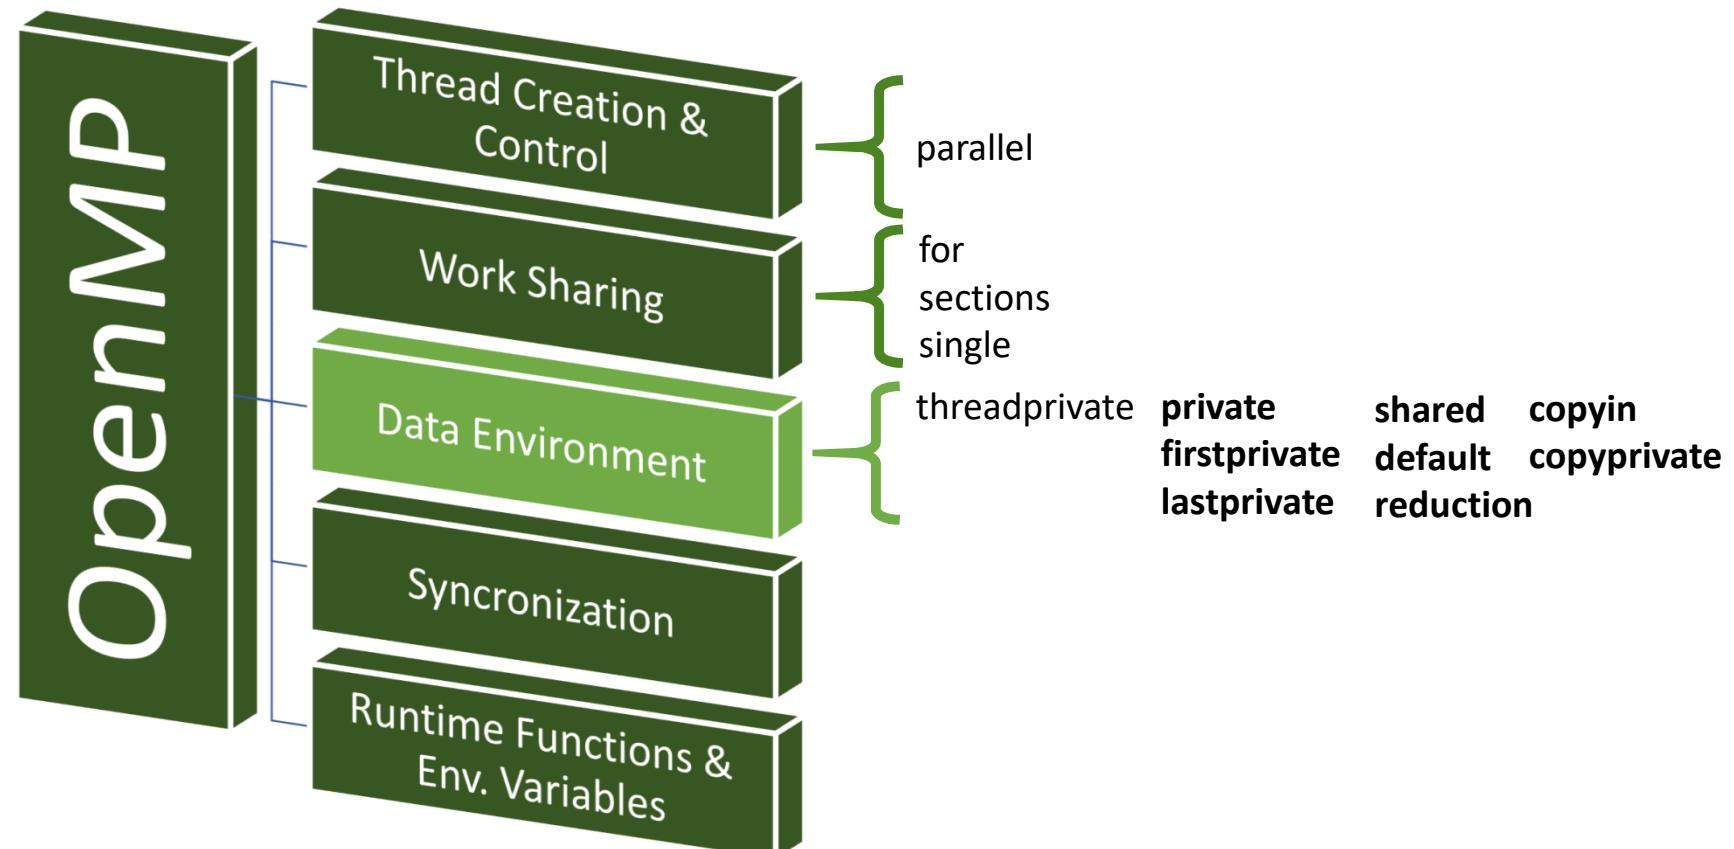


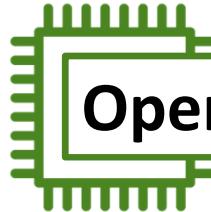
## OpenMP – NOTES on Work Sharing Constructs



All worksharing constructs have an implied barrier at the end.

# OpenMP Core Elements





## OpenMP – Data Environment

### *Important notes on variables:*

- If a variable is visible when a parallel or work-sharing construct is encountered, and the variable is not specified in a sharing attribute clause or threadprivate directive, then the variable is shared.
- Static variables declared within the dynamic extent of a parallel region are shared.
- Heap allocated memory (for example, using malloc() in C or C++ or the new operator in C++) is shared.
- The pointer to previous point's memory, however, can be either private or shared.

## OpenMP – Data Environment

### **threadprivate** directive

- The **threadprivate** directive specifies that variables are replicated, with each thread having its own copy.

```
#pragma omp threadprivate(variable-list) new-line
```



- A **private** variable is local to a region and will be placed on the stack.
- A **threadprivate** variable on the other hand will be most likely placed in the heap.
- Each copy of an **omp threadprivate** data variable is initialized once prior to first use of that copy.
- A threadprivate variable persist across regions (depending on some restrictions).

# OpenMP – Data Environment

## **private** clause

- The **private** clause declares the variables in *variable-list* to be private to each thread in a team.

```
private (variable-list)
```

# OpenMP – Data Environment

## ***firstprivate*** clause

- The ***firstprivate*** clause provides a superset of the functionality provided by the ***private*** clause.

```
firstprivate(variable-list)
```

- The main difference versus the ***private*** clause is that the value of the variable is the value of the original object that exists immediately prior to the parallel or the work-sharing construct.

# OpenMP – Data Environment

## ***lastprivate*** clause

- The ***lastprivate*** clause provides a superset of the functionality provided by the ***private*** clause.

```
lastprivate (variable-list)
```

- When a ***lastprivate*** clause appears on the directive that identifies a work-sharing construct, the value of each ***lastprivate*** variable from the sequentially last iteration of the associated loop, or the lexically last section directive, is assigned to the variable's original object.

# OpenMP – Data Environment

## **shared** clause

- This clause shares variables that appear in the *variable-list* among all the threads in a team.
- All threads within a team access the same storage area for **shared** variables.

**shared**(*variable-list*)

# OpenMP – Data Environment

## **default** clause

- The **default** clause allows the user to affect the data-sharing attributes of variables.

```
default (shared | none)
```

- Specifying **default (shared)** is equivalent to explicitly listing each currently visible variable in a **shared** clause, unless it is **threadprivate** or **const** qualified.
- In the absence of an explicit **default** clause, the default behavior is the same as if **default (shared)** were specified.

# OpenMP – Data Environment

## *reduction* clause

- This clause performs a reduction on the scalar variables that appear in *variable-list*, with the operator *op*.

### reduction (*op*:*variable-list*)

```
x = x op expr
x binop= expr
x = expr op x      (except for subtraction)
x++
++x
x--
--x
```

#### where:

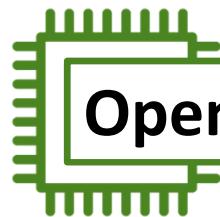
<i>x</i>	One of the reduction variables specified in the <i>list</i> .
<i>variable-list</i>	A comma-separated list of scalar reduction variables.
<i>expr</i>	An expression with scalar type that does not reference <i>x</i> .
<i>op</i>	Not an overloaded operator but one of +, *, -, &, ^,  , &&, or   .
<i>binop</i>	Not an overloaded operator but one of +, *, -, &, ^, or  .

## OpenMP – Data Environment

### ***copyin*** clause

- The ***copyin*** clause provides a mechanism to assign the same value to **threadprivate** variables for each thread in the team executing the parallel region.
- For each variable specified in a ***copyin*** clause, the value of the variable in the master thread of the team is copied, as if by assignment, to the thread-private copies at the beginning of the parallel region.

```
copyin (variable-list)
```



## OpenMP – Data Environment

### **copyprivate** clause

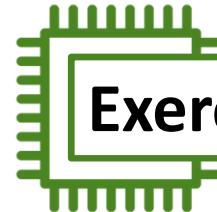
- The **copyprivate** clause provides a mechanism to use a private variable to broadcast a value from one member of a team to the other members.
- It is an alternative to using a shared variable for the value when providing such a shared variable would be difficult (for example, in a recursion requiring a different variable at each level).

```
copyprivate (variable-list)
```



- The **copyprivate** clause can only appear on the **single** directive.

threadprivate  
private  
firstprivate  
lastprivate  
shared  
default  
reduction  
copyin  
copyprivate



## Exercise: Testing OpenMP... Data Environment

The global variable **var** has a race condition. Use clauses to correct the condition. Each thread wants its own copy of **var**, so you should try to make their different values persist from one parallel region to the other.

```
#include <omp.h>
#include <stdio.h>
int var = 10;

void main()
{
    printf("Initial value (%d), Var=%d\n\n", omp_get_thread_num(), var);
    #pragma omp parallel num_threads (2)
    {
        var *= (omp_get_thread_num() + 1); //Potential race condition.
        printf("Thread (%d), Var=%d\n\n", omp_get_thread_num(), var);
    }
    var += 1000;
    #pragma omp parallel num_threads (2)
    {
        printf("Thread (%d), Var=%d\n\n", omp_get_thread_num(), var);
    }
}
```

```
Initial value (0), Var=10
Thread (0), Var=10
Thread (1), Var=20
Thread (0), Var=1020
Thread (1), Var=1020
```



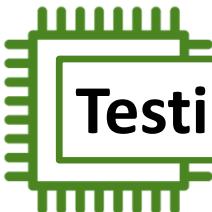
## Testing OpenMP... Data Environment

```
#include <omp.h>
#include <stdio.h>

int var = 10;
#pragma omp threadprivate (var) //Only valid for global or static data items

void main()
{
    printf("Initial value (%d), Var=%d\n\n", omp_get_thread_num(), var);
    #pragma omp parallel num_threads (2)
    {
        var *= (omp_get_thread_num() + 1); //Potential race condition.
        printf("Thread (%d), Var=%d\n\n", omp_get_thread_num(), var);
    }
    var += 1000;
    #pragma omp parallel num_threads (2)
    {
        printf("Thread (%d), Var=%d\n\n", omp_get_thread_num(), var);
    }
}
```

```
Initial value (0), Var=10
Thread (0), Var=10
Thread (1), Var=20
Thread (0), Var=1010
Thread (1), Var=20
```



threadprivate  
private  
firstprivate  
lastprivate  
shared  
default  
reduction  
copyin  
copyprivate

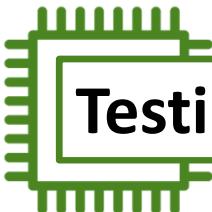
## Testing OpenMP... Data Environment

In the second parallel region, I want to have the same value of the variable var. How can I do that?

```
#include <omp.h>
#include <stdio.h>

int var = 10;
#pragma omp threadprivate (var) //Only valid for global or static data items

void main()
{
    printf("Initial value (%d), Var=%d\n\n", omp_get_thread_num(), var);
    #pragma omp parallel num_threads (2)
    {
        var *= (omp_get_thread_num() + 1); //Potential race condition.
        printf("Thread (%d), Var=%d\n\n", omp_get_thread_num(), var);
    }
    var += 1000;
    #pragma omp parallel num_threads (2)
    {
        printf("Thread (%d), Var=%d\n\n", omp_get_thread_num(), var);
    }
}
```



threadprivate  
private  
firstprivate  
lastprivate  
shared  
default  
reduction  
copyin  
copyprivate

## Testing OpenMP... Data Environment

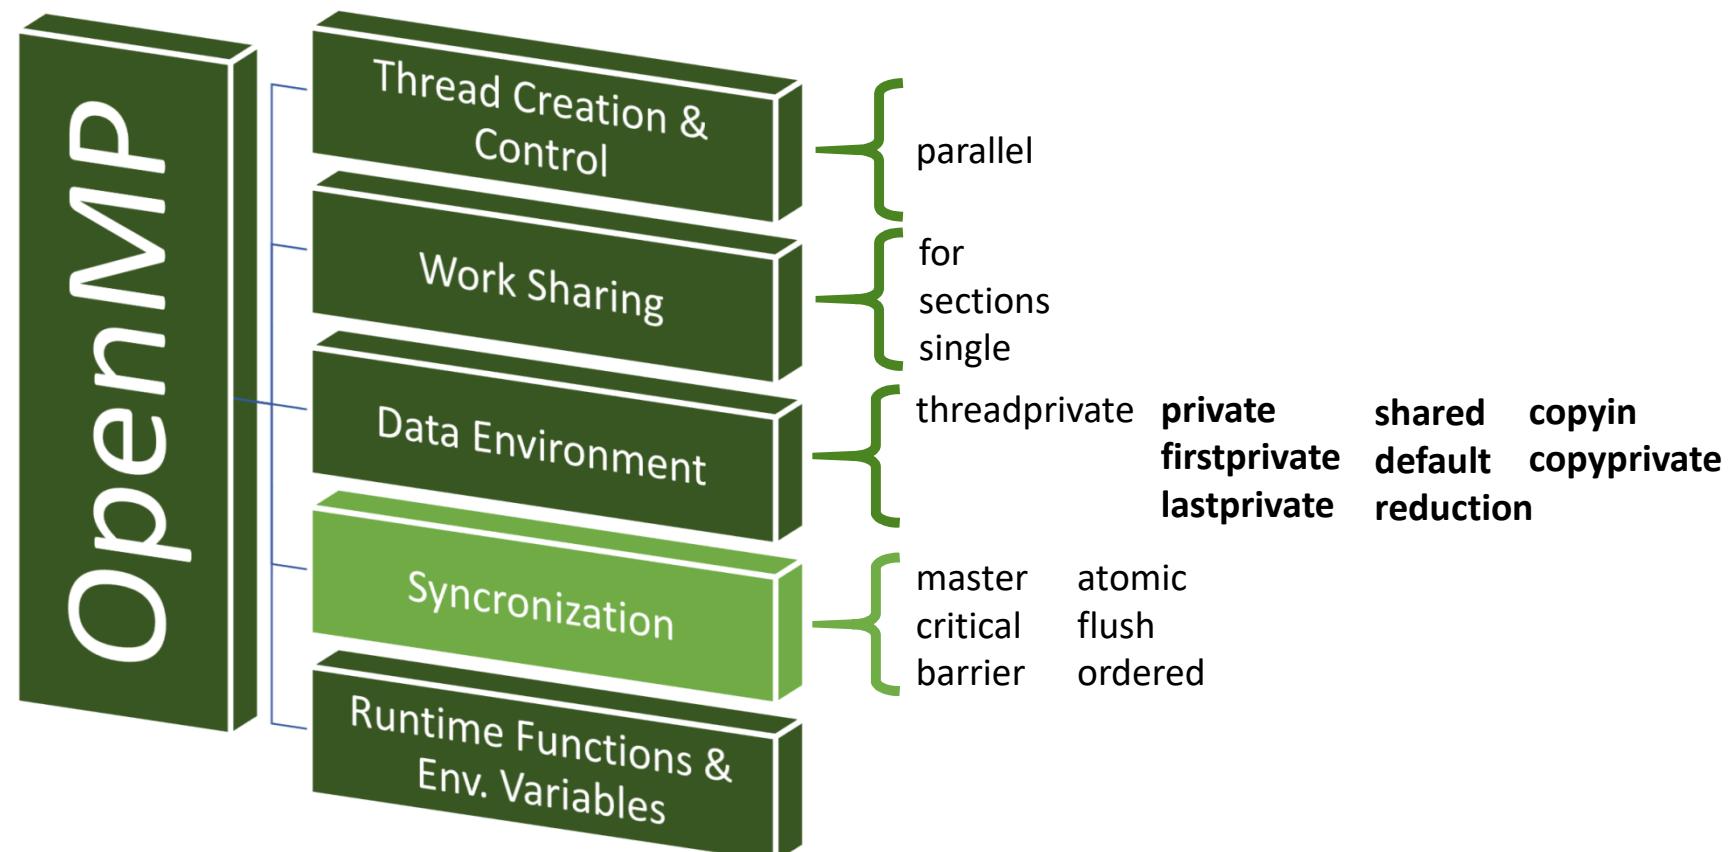
In the second parallel region, I want to have the same value of the variable var. How can I do that?

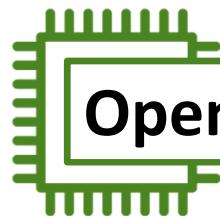
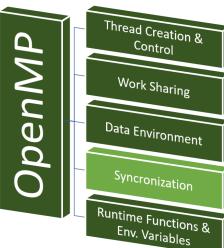
```
#include <omp.h>
#include <stdio.h>

int var = 10;
#pragma omp threadprivate (var) //Only valid for global or static data items

void main()
{
    printf("Initial value (%d), Var=%d\n\n", omp_get_thread_num(), var);
    #pragma omp parallel num_threads (2)
    {
        var *= (omp_get_thread_num() + 1); //Potential race condition.
        printf("Thread (%d), Var=%d\n\n", omp_get_thread_num(), var);
    }
    var += 1000;
    #pragma omp parallel copyin (var) num_threads (2)
    {
        printf("Thread (%d), Var=%d\n\n", omp_get_thread_num(), var);
    }
}
```

# OpenMP Core Elements





## OpenMP – Master and Synchronization

### **master** construct

- The **master** directive identifies a construct that specifies a structured block that is executed by the master thread of the team.

```
#pragma omp master new-line  
structured-block
```



- There is no implied barrier either on entry to or exit from the master construct.

# OpenMP – Master and Synchronization

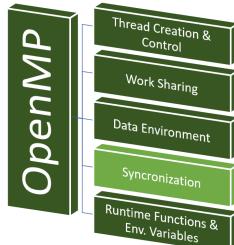
## ***critical*** construct

- The ***critical*** directive identifies a construct that restricts execution of the associated structured block to a single thread at a time.

```
#pragma omp critical [(name)] new-line
structured-block
```



- An optional *name* may be used to identify the critical region. Identifiers used to identify a critical region have external linkage and are in a name space which is separate from the name spaces used by labels, tags, members, and ordinary identifiers.
- A thread waits at the beginning of a critical region until no other thread is executing a critical region (anywhere in the program) with the same name. All unnamed ***critical*** directives map to the same unspecified name.



# OpenMP – Master and Synchronization

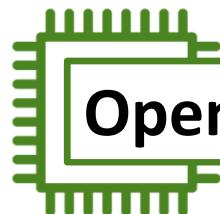
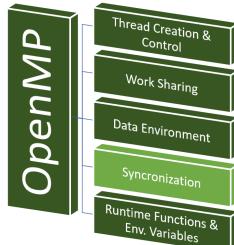
## **barrier** construct

- The **barrier** directive synchronizes all the threads in a team. When encountered, each thread in the team waits until all of the others have reached this point.

```
#pragma omp barrier new-line
```



- After all threads in the team have encountered the barrier, each thread in the team begins executing the statements after the barrier directive in parallel

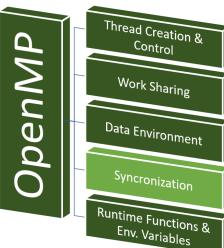


## OpenMP – Master and Synchronization

### **atomic** construct

- The **atomic** directive ensures that a specific memory location is updated atomically, rather than exposing it to the possibility of multiple, simultaneous writing threads.

```
#pragma omp atomic new-line  
expression-stmt
```



# OpenMP – Master and Synchronization

## **flush** construct

- The **flush** directive specifies a “cross-thread” sequence point at which the implementation is required to ensure that all threads in a team have a consistent view of certain objects (specified below) in memory.
- Explicit form:

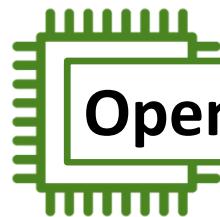
```
#pragma omp flush [(variable-list)] new-line
```
- A flush directive without a variable-list synchronizes all shared objects except inaccessible objects with automatic storage duration.

## OpenMP – Master and Synchronization

### *flush* construct

- In can also be implied (with no variable list):
  - ❖ At **barrier**
  - ❖ At entry to and exit from **critical**
  - ❖ At entry to and exit from **ordered**
  - ❖ At entry to and exit from **parallel**
  - ❖ At exit from **for**
  - ❖ At exit from **sections**
  - ❖ At exit from **single**
  - ❖ At entry to and exit from **parallel for**
  - ❖ At entry to and exit from **parallel sections**





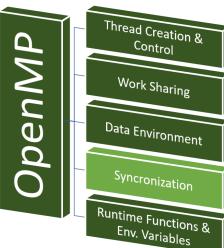
## OpenMP – Master and Synchronization

### **flush** construct

➤ **flush** directive is not implied for any of the following:

- ❖ If a **nowait** clause is present.
- ❖ At entry to **for**
- ❖ At entry to or exit from **master**
- ❖ At entry to **sections**
- ❖ At entry to **single**





# OpenMP – Master and Synchronization

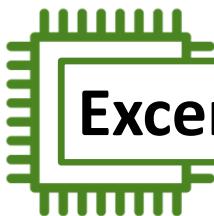
## *ordered* construct

- The structured block following an **ordered** directive is executed in the order in which iterations would be executed in a sequential loop.

```
#pragma omp ordered new-line  
structured-block
```



- The **ordered** construct sequentializes and orders the execution of ordered regions while allowing code outside the region to run in parallel.



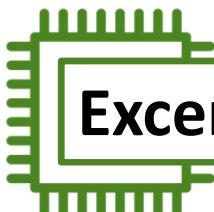
## Excercise - #pragma omp ordered

Use directive **#pragma omp ordered** and clause **ordered** to print the results in the order they would be printed if executed sequentially.

```
#include <stdio.h>

void main()
{
    int i;
    #pragma omp parallel for num_threads(4)
    for (i = 0; i < 20; i++)
    {
        printf(" %d\n", i);
    }
}
```

```
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
0
1
2
3
4
```



## Excercise - #pragma omp ordered

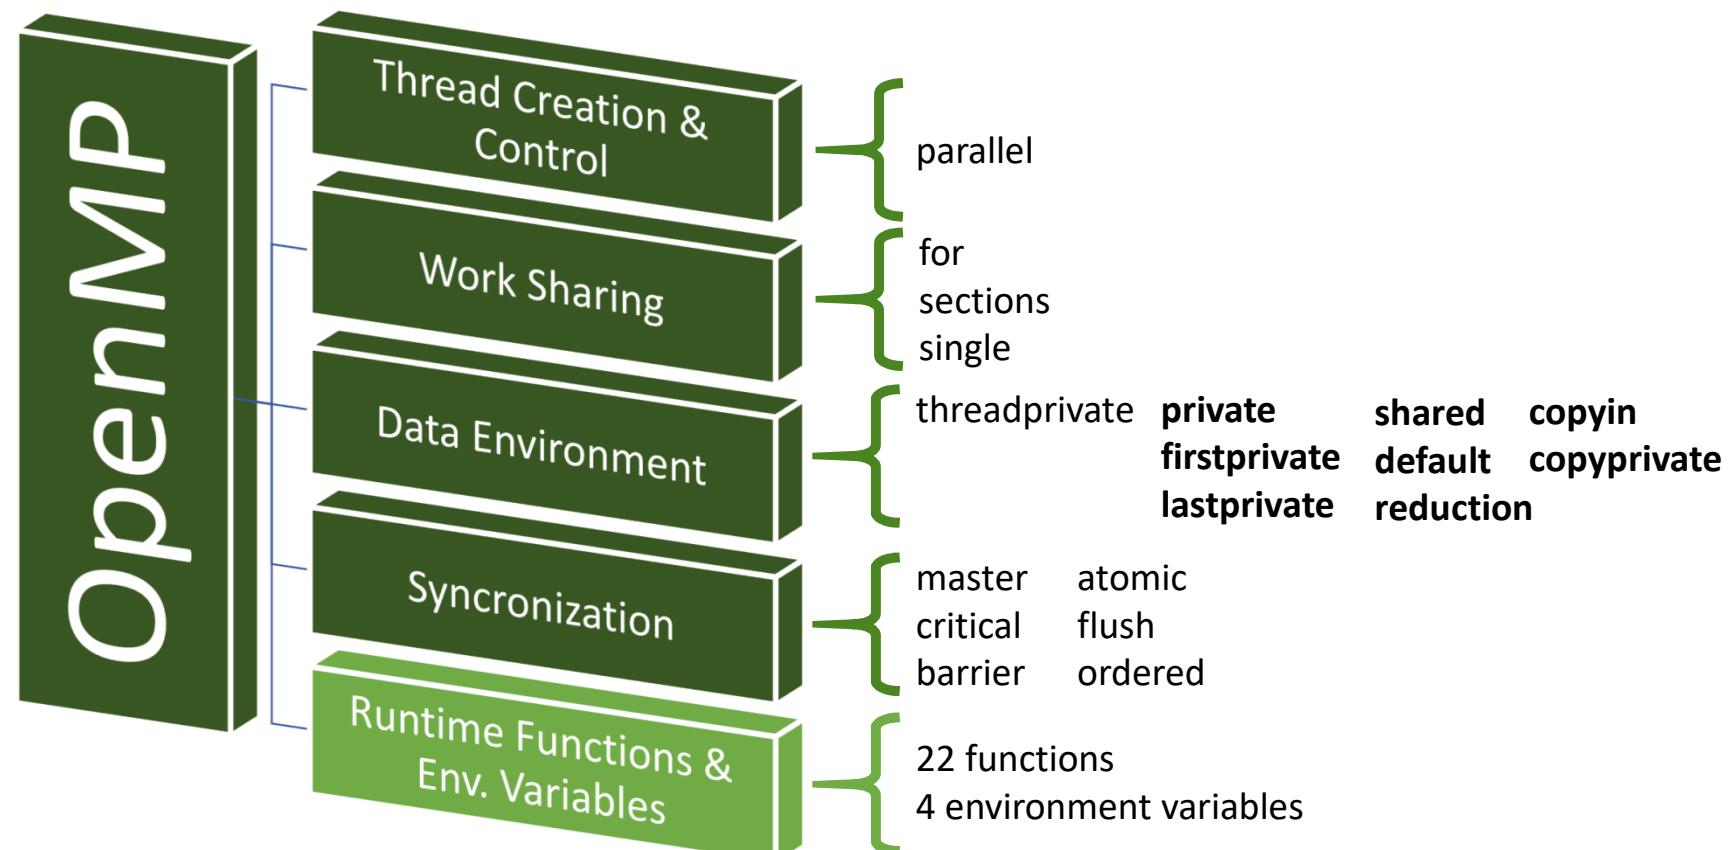
Use directive `#pragma omp ordered` and clause `ordered` to print the results in the order they would be printed if executed sequentially.

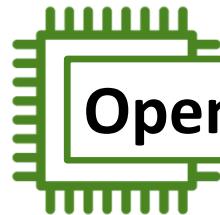
```
#include <stdio.h>

void main()
{
    int i;
    #pragma omp parallel for ordered num_threads(4)
    for (i = 0; i < 20; i++)
    {
        #pragma omp ordered
        printf(" %d\n", i);
    }
}
```

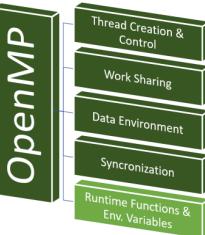
```
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
```

# OpenMP Core Elements

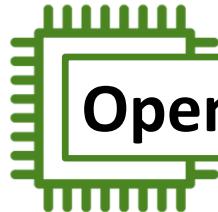




# OpenMP – Functions and Environmental Variables

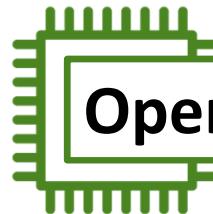


Function	
Execution Environment	



# OpenMP – Functions and Environmental Variables

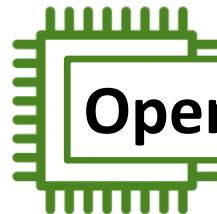
	Function	
Execution Environment	omp_set_num_threads	sets the default number of threads to use for subsequent parallel regions
	omp_get_num_threads	returns the number of threads currently in the team executing the parallel region from which it is called.
	omp_get_max_threads	returns an integer that is guaranteed to be at least as large as the number of threads that would be used to form a team if a parallel region without a num_threads clause were to be encountered at that point in the code.
	omp_get_thread_num	returns the thread number, within its team, of the thread executing the function. The thread number lies between 0 and omp_get_num_threads()–1, inclusive. The master thread of the team is thread 0.
	omp_get_num_procs	returns the number of processors that are available to the program at the time the function is called.
	omp_in_parallel	returns a nonzero value if it is called within the dynamic extent of a parallel region executing in parallel; otherwise, it returns 0.
	omp_set_dynamic	enables or disables dynamic adjustment of the number of threads available for execution of parallel regions.
	omp_get_dynamic	returns a nonzero value if dynamic adjustment of threads is enabled, and returns 0 otherwise.
	omp_set_nested	enables or disables nested parallelism.
	omp_get_nested	returns a nonzero value if nested parallelism is enabled and 0 if it is disabled.



# OpenMP – Functions and Environmental Variables

	Function	Nestable Function	Description
Lock			

	Function	Description
Timing		



# OpenMP – Functions and Environmental Variables

	<b>Function</b>	<b>Nestable Function</b>	<b>Description</b>
Lock	omp_init_lock	omp_init_nest_lock	initializes a simple/nestable lock.
	omp_destroy_lock	omp_destroy_nest_lock	removes a simple/nestable lock.
	omp_set_lock	omp_set_nest_lock	waits until a simple/nestable lock is available.
	omp_unset_lock	omp_unset_nest_lock	releases a simple/nestable lock.
	omp_test_lock	omp_test_nest_lock	tests a simple/nestable lock.

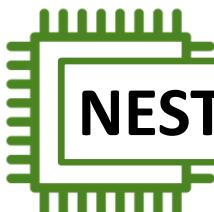
	<b>Function</b>	<b>Description</b>
Timing	omp_get_wtime	returns elapsed wall-clock time.
	omp_get_wtick	returns seconds between successive clock ticks.

# OpenMP – Functions and Environmental Variables

## Environmental Variable

# OpenMP – Functions and Environmental Variables

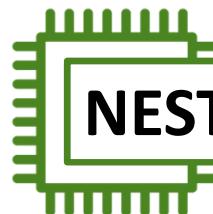
Environmental Variable		
OMP_SCHEDULE	sets the run-time schedule type and chunk size.	static,[chunk size] dynamic,[chunk size] guided,[chunk size]
OMP_NUM_THREADS	sets the number of threads to use during execution.	Positive integer
OMP_DYNAMIC	enables or disables dynamic adjustment of the number of threads.	TRUE   FALSE
OMP_NESTED	enables or disables nested parallelism.	TRUE   FALSE (default)



Normally **DISABLED** by default. You can enable it using a function or an environment variable.

```
#include <omp.h>
#include <stdio.h>

void main()
{
    //omp_set_nested(1);
    #pragma omp parallel num_threads(2)
    {
        printf("Level 1: number of threads in the team - %d\n", omp_get_num_threads());
        #pragma omp parallel num_threads(2)
        {
            printf("Level 2: number of threads in the team - %d\n", omp_get_num_threads());
            #pragma omp parallel num_threads(2)
            {
                printf("Level 3: number of threads in the team - %d\n", omp_get_num_threads());
            }
        }
    }
}
```



# NESTED

Nested disabled...

CS Microsoft Visual Studio Debug Console

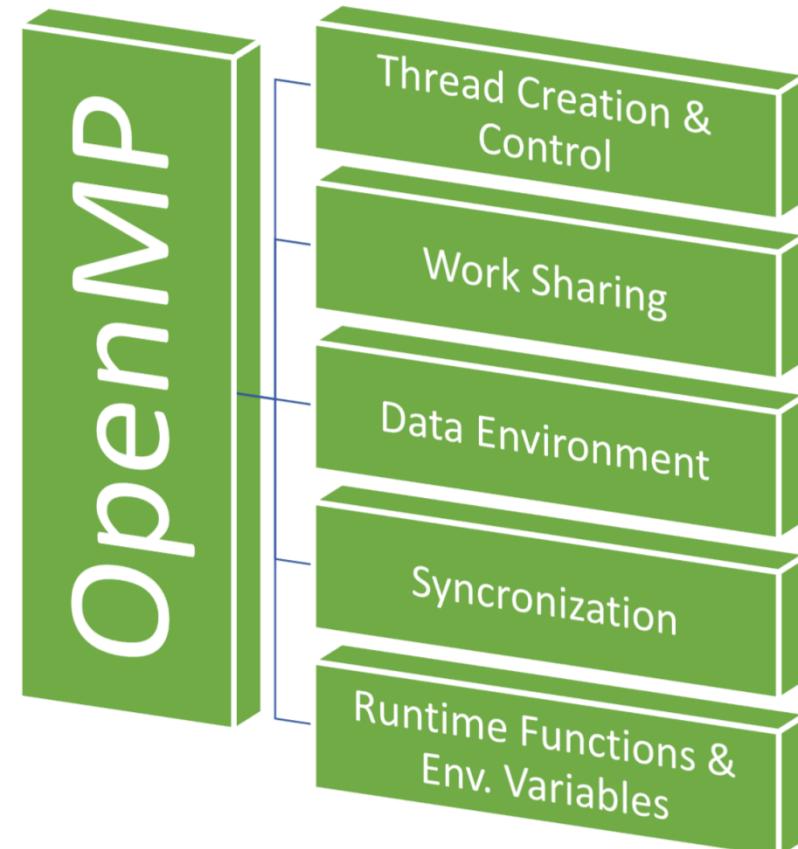
```
Level 1: number of threads in the team - 2  
Level 2: number of threads in the team - 1  
Level 3: number of threads in the team - 1  
Level 1: number of threads in the team - 2  
Level 2: number of threads in the team - 1  
Level 3: number of threads in the team - 1
```

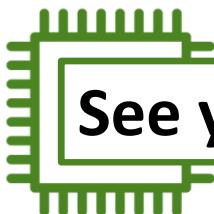
Nested enabled...

Microsoft Visual Studio Debug Console

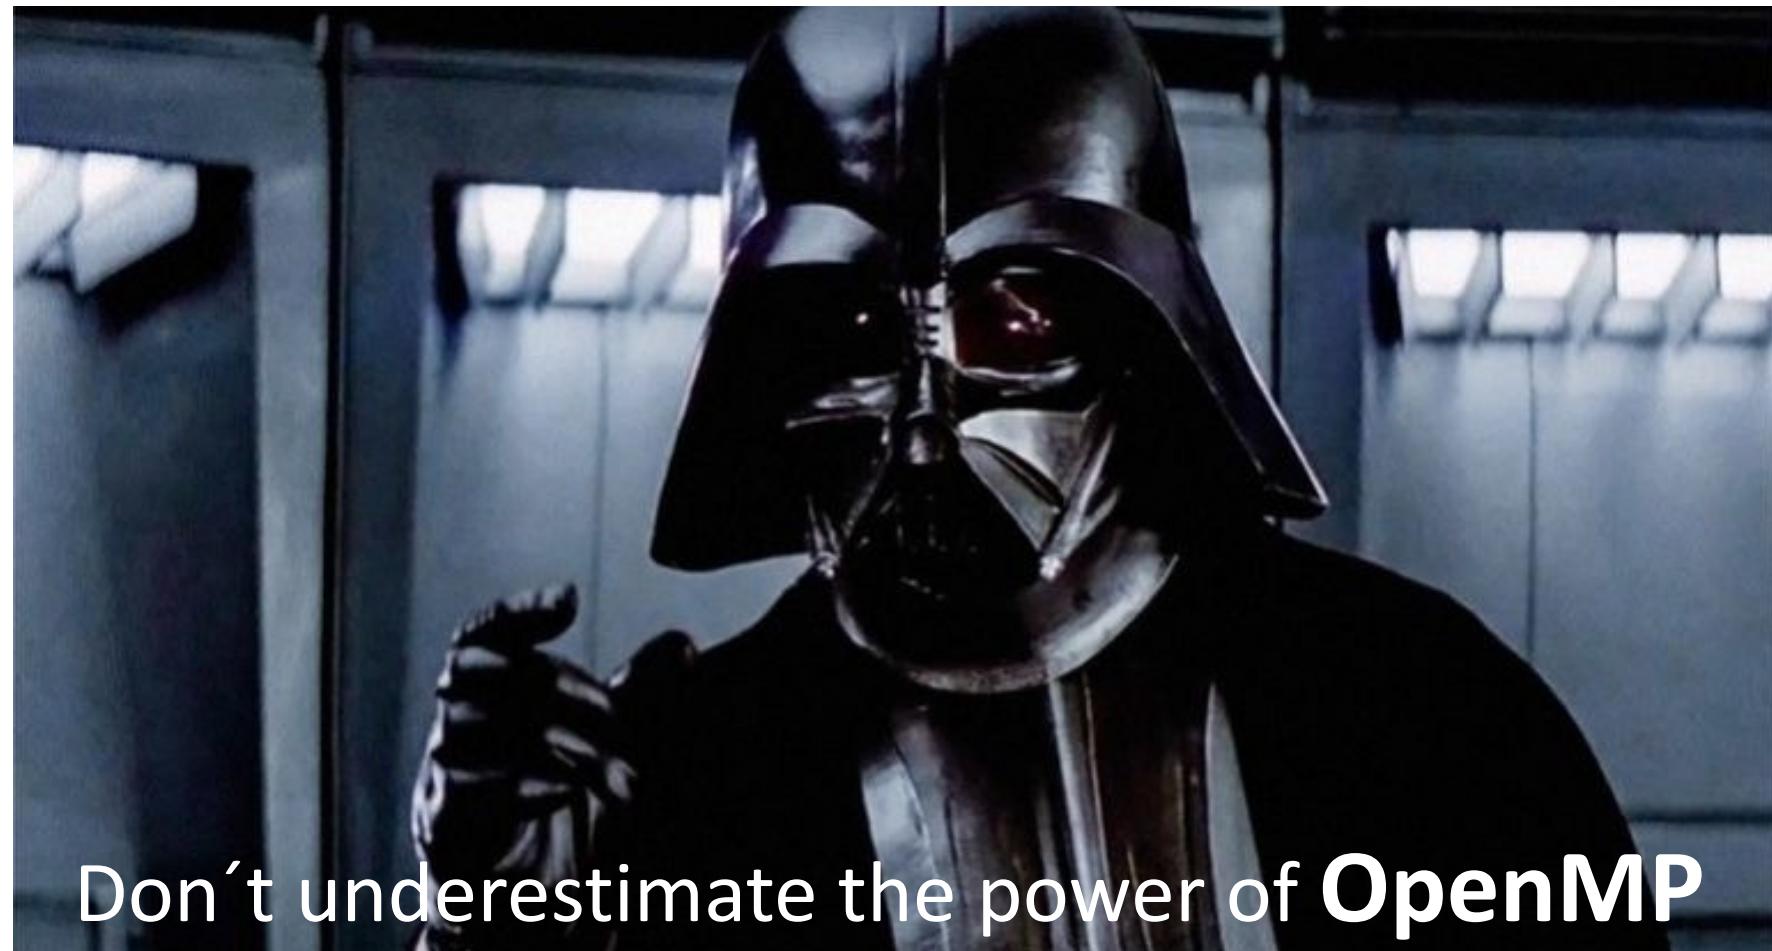
```
Level 1: number of threads in the team - 2
Level 1: number of threads in the team - 2
Level 2: number of threads in the team - 2
Level 2: number of threads in the team - 2
Level 3: number of threads in the team - 2
Level 2: number of threads in the team - 2
Level 3: number of threads in the team - 2
Level 3: number of threads in the team - 2
Level 3: number of threads in the team - 2
Level 2: number of threads in the team - 2
Level 3: number of threads in the team - 2
Level 3: number of threads in the team - 2
Level 3: number of threads in the team - 2
Level 3: number of threads in the team - 2
```

# OpenMP – Questions





See you soon...



Don't underestimate the power of OpenMP