



Multiprocessors

Feb-Jun 2021

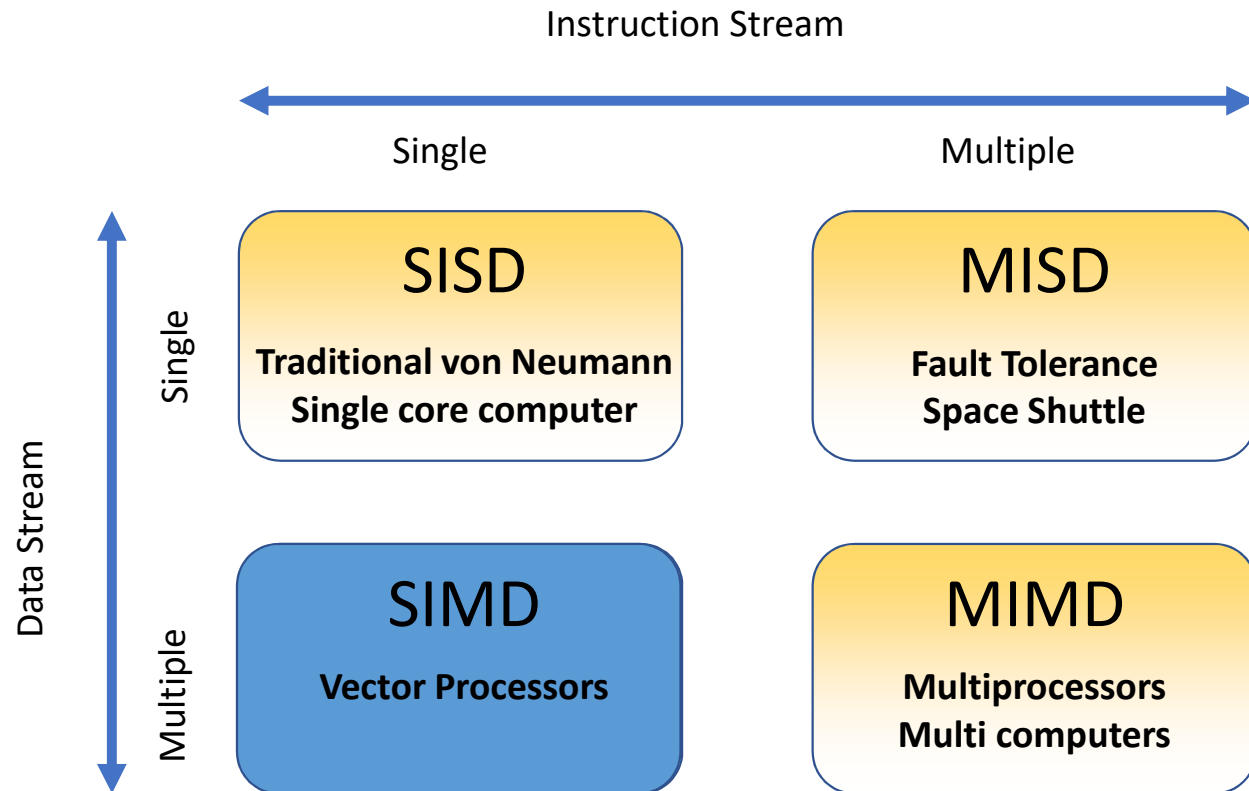
Vectorization

2

Alejandro Guajardo Moreno



Flynn's Taxonomy



Vectorization Foucs: SIMD



SIMD in the real world?



Can you think of an example where SIMD can be used?

EXAMPLE

Hint: SAME Instruction over lots and lots of Data.

- ☐ Image Processing
- ☐ 3D graphics
- ☐ Rendering
- ☐ Criptography
- ☐ Compression
- ☐ Video Processing
 - ☐ conversion between video standards and frame rates (NTSC to/from PAL, NTSC to/from HDTV formats)
 - ☐ deinterlacing
 - ☐ image noise reduction
 - ☐ adaptive video compression
 - ☐ image enhancement
- ☐ And many many MORE!

Vectorization comparison with mundane objects



Vectorization!



What is VECTORIZATION and its relation to SIMD?

Scalar Instructions

$$\begin{array}{rcl} 4 & + & 1 = 5 \\ 0 & + & 3 = 3 \\ -2 & + & 8 = 6 \\ 9 & + & -7 = 2 \end{array}$$



Vector Instructions

$$\begin{array}{rcl} 4 & & 1 & = & 5 \\ 0 & & 3 & = & 3 \\ -2 & + & 8 & = & 6 \\ 9 & & -7 & = & 2 \end{array} \quad \begin{array}{c} \updownarrow \\ \text{Vector Length} \end{array}$$

Example... Scalar instructions...

1. Load (A)
2. Load (B)
3. C = Sum (A+B)
4. Store (C)

32 bit signed integer





Scalar Instructions

4	+	1	=	5
0	+	3	=	3
-2	+	8	=	6
9	+	-7	=	2

Example... Scalar instructions...

1. Load (A)
 2. Load (B)
 3. C = Sum (A+B)
 4. Store (C)
- 5 - 8. Repeat with $0 + 3 = 3$
9 - 12. Repeat with $-2 + 8 = 6$
13 - 16. Repeat with $9 + -7 = 2$

32 bit signed integer



Scalar Instructions

4	+	1	=	5
0	+	3	=	3
-2	+	8	=	6
9	+	-7	=	2

Example... Now using **VECTOR** instructions...


Vector Instructions

4	1	5
0	3	3
-2	8	6
9	-7	2


+ =

Vector Length

32 bit signed integer



1. Load (vA)
2. Load (vB)
3. vC = Sum (vA+vB)
4. Store (vC)



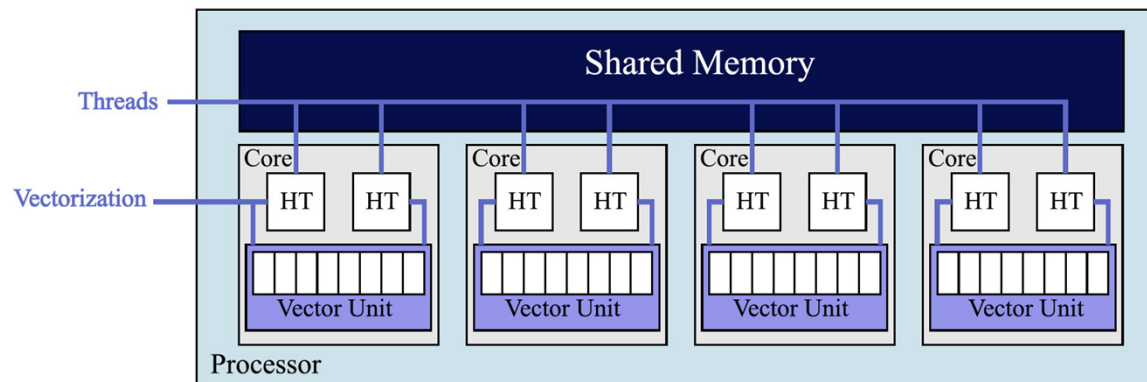
What did you notice?

Memory?

- Elements are located next to each other. **Contiguous.**
- Elements are normally and preferable **aligned.**

Registers?

- “Special Looooong Registers”





Components needed for a processor to use vectors.

- With VECTOR Instruction Set Extensions.
- With Special Registers, transistors & logic.



First INTEL® VECTOR Instructions



- 1996
- CISC
- Superscalar
- IA-32 ← 32-bit version of the x86 instruction set architecture

➤ MMX = SIMD (64 bit)

MMX = Multimedia Extensions


And after...

MMX: Multimedia Extension

SSE: Streaming SIMD Extension

AVX: Advanced Vector Extension

x86-16	8086
	286
x86-32	386
	486
MMX	Pentium
	Pentium MMX
SSE	Pentium III
	Pentium 4
SSE2	Pentium 4E
	Pentium 4F
SSE3	Core 2 Duo
	Penryn
SSE4	Core i7 (Nehalem)
	Sandy Bridge
AVX	Haswell
	Skylake X
AVX2	
AVX-512	



And after...

Vendor	Name		ν -way	Precision	Introduced with
Intel	SSE		4-way	single	Pentium III
	SSE2	+	2-way	double	Pentium 4
	SSE3				Pentium 4 (Prescott)
	SSSE3				Core Duo
	SSE4				Core2 Extreme (Penryn)
	AVX		8-way 4-way	single double	Core i7 (Sandybridge)
Intel	IPF		2-way	single	Itanium
Intel	LRB		16-way	single	Larrabee
			8-way	double	
AMD	3DNow!		2-way	single	K6
	Enhanced 3DNow!				K7
	3DNow! Professional	+	4-way	single	Athlon XP
	AMD64	+	2-way	double	Opteron
Motorola	AltiVec		4-way	single	MPC 7400 G4
IBM	VMX		4-way	single	PowerPC 970 G5
	SPU	+	2-way	double	Cell BE
IBM	Double FPU		2-way	double	PowerPC 440 FP2

Source: ([CS263-2300 ETH](#))



Teams

3

Use <https://ark.intel.com> or <https://www.amd.com/en/products/specifications> to answer the following:

- ❖ What SIMD extensions does the i5-8350U processor support?

Intel® SSE4.1, Intel® SSE4.2, Intel® AVX2

- ❖ What is the targeted segment for i5-8350U processor?

Mobile

- ❖ Does the Intel Itanium Processor 9760 has any kind of SIMD extensions?

No extensions*

- ❖ Intel Xeon Platinum 9282, second generation of Intel Xeon Scalable Processors support for SIMD extensions?

Intel® AVX-512

- ❖ Is there an Intel® ATOM processor that supports AVX, AV2 or AVX-512?

No

- ❖ AMD Ryzen™ Threadripper™ 1950X Processor supports AVX2 and/or AVX-512?

AVX2

IA-32 Architecture

- Datatypes: BYTE, WORD, DWORD

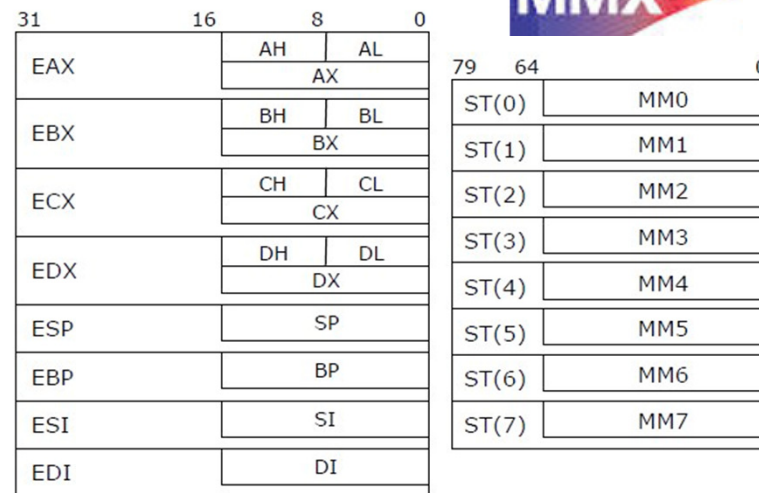

31	16	8	0
EAX	AH	AL	
	AX		
EBX	BH	BL	
	BX		
ECX	CH	CL	
	CX		
EDX	DH	DL	
	DX		
ESP	SP		
EBP	BP		
ESI	SI		
EDI	DI		

General Purpose Registers

x86-16	8086
	286
	386
	486
	Pentium
	Pentium MMX
	SSE
	Pentium III
	SSE2
	Pentium 4
	SSE3
	Pentium 4E
	Pentium 4F
x86-64/em64t	
	Core 2 Duo
	SSE4
	Penryn
	Core i7 (Nahalem)

IA-32 Architecture

- Datatypes: BYTE, WORD, DWORD y **QWORD**
- The ST(x) registers are accessible as a stack by the FPU, while the MMX registers are accessible by specific instructions.
- For extended precision (80 bits) floating point operations the vectorization is impossible.

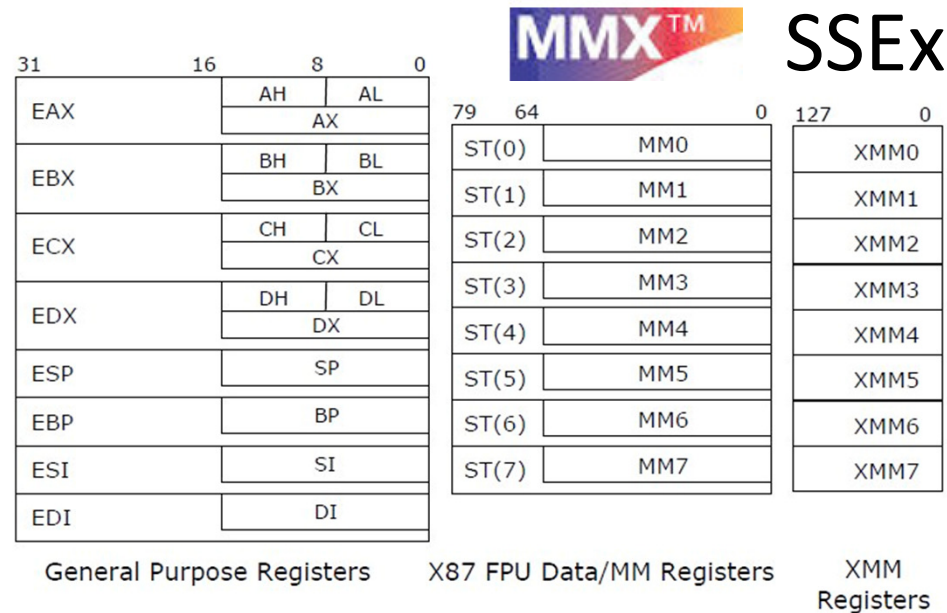



General Purpose Registers

X87 FPU Data/MM Registers

IA-32 Architecture

- Datatypes: BYTE, WORD, DWORD y QWORD
- The ST(x) registers are accessible as a stack by the FPU, while the MMX registers are accessible by specific instructions.
- For extended precision (80 bits) floating point operations the vectorization is impossible.





IA-32 Architecture

When data is vectorial, the processor may use any of two options:

- To use the **FPU (x87)**
- To use **XMM** registers

```
float x[4], y[4], z[4];
int i;
...
for (i=0; i<4; i++)
    z[i]=x[i]*y[i];
```

Assembly for Intel SSE

```
movaps  x, %xmm0
mulps   y, %xmm0
movaps  %xmm0, z
```

Assembly for x87

```
flds    x
fmuls   y
flds    4+x
fmuls   4+y
flds    8+x
fmuls   8+y
flds    12+x
fmuls   12+y
fxch    %st(3)
fstps   z
fxch    %st(1)
fstps   4+z
fstps   8+z
fstps   12+z
```

Which one is better in
terms of
performance/IPC?

x86-64 Architecture

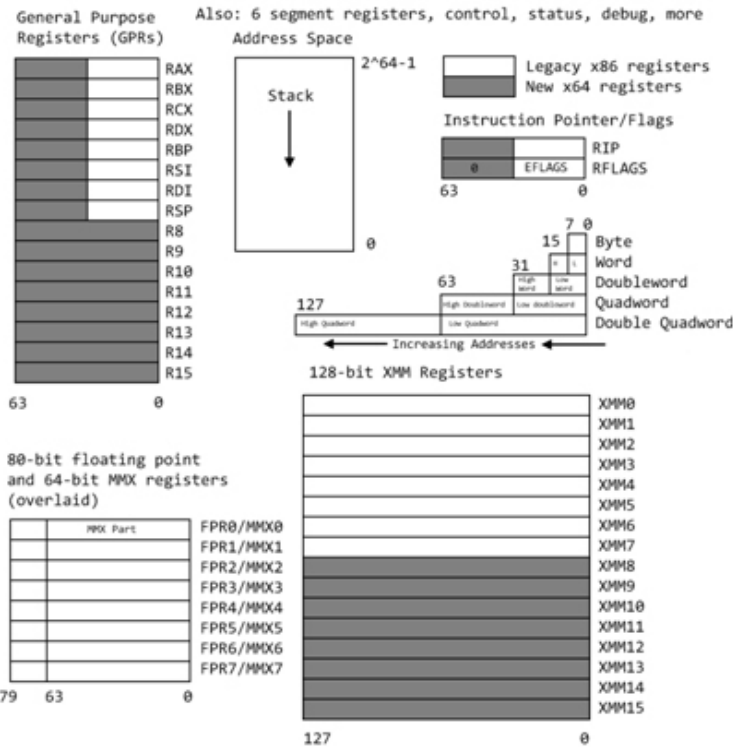
63	32	16	8	0
RAX	EAX	AH	AL	
		AX		
RBX	EBX	BH	BL	
		BX		
RCX	ECX	CH	CL	
		CX		
RDX	EDX	DH	DL	
		DX		
RSP	ESP	SP	SPL	
RBP	EBP	BP	BPL	
RSI	ESI	SI	SIL	
RDI	EDI	DI	DIL	
R8	R8D	R8W	R8B	
R9	R9D	R9W	R9B	
...				
R15	R15D	R15W	R15B	

General Purpose Registers


127	0
XMM0	
XMM1	
XMM2	
XMM3	
...	
XMM15	

XMM Registers

Intel 64 / AMD x86-64 Full Architectures



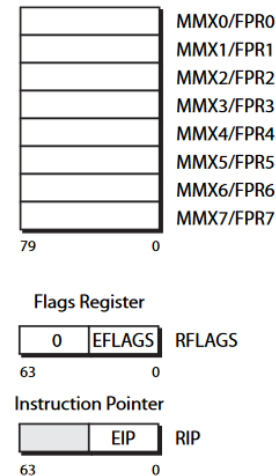
General-Purpose Registers (GPRs)




Legacy x86 registers, supported in all modes

Register extensions, supported in 64-bit mode

64-Bit Media and Floating-Point Registers



SSE Media Registers




Application-programming registers not shown include Media eXtension Control and Status Register (MXCSR) and x87 tag-word, control-word, and status-word registers

513-101 ymmleaps






SSE Family




- ✓ Each version incorporate new instructions.
- ✓ Backward compatibility is guaranteed.
- ✓ Each core can execute SSE instructions independently.

Data extensions (SSE)

- Floating point vectors
 - 4-way single precision (DWORD)




- Floating point scalars
 - Single precision (DWORD)



Data extensions (SSE2)


- Integer vectors

- 16-way bytes (BYTE)
- 8-way 2 bytes (WORD)
- 4-way 4 bytes (DWORD)
- 2-way 8 bytes (QWORD)




- Floating point vectors

- 4-way single precision (DWORD)
- 2-way double precision (QWORD)



- Floating point scalars

- Single precision (DWORD)
- Double precision (QWORD)





Data extensions (SSE3)

Focus:

Horizontal work within the register.

HADDPD — (Horizontal-Add-Packed-Double)

HADDPS (Horizontal-Add-Packed-Single)

HSUBPD — (Horizontal-Subtract-Packed-Double)

HSUBPS — (Horizontal-Subtract-Packed-Single)

Also:

3D and Digital Signal Processing operations.

Convert FP to Int avoiding pipeline stall.



Data extensions (SSSE3)

Focus:

Integers

Use of either MMX or XMM registers.



Data extensions (SSE4)

Focus:

Non multimedia.

Dot product.

Absolute, Max/Min, Round

String and Text New Instructions



Advanced Vector eXtensions - AVX

- Width of registers increased from 128 to 256 bits.
- 8 registers in 32 bit mode: YMM0 – YMM7.
- 16 registers in 64 bit mode: YMM0 – YMM15.
- Three-operand SIMD instruction format vs two-operand in SSE.
- The alignment requirement of SIMD memory operands is relaxed.



Advanced Vector eXtensions 2 – AVX2

- Three-operand general-purpose bit manipulation and multiply
- Gather support, enabling vector elements to be loaded from non-contiguous memory locations
- DWORD- and QWORD-granularity any-to-any permutes
- Vector shifts.
- Three-operand fused multiply-accumulate support (**FMA3**)

$$a = a \cdot c + b$$

$$a = b \cdot a + c$$

$$a = b \cdot c + a$$



Note on FMA4...

- Four-operand fused multiply-accumulate support (FMA4)

$$a = b \cdot c + d$$

Only available in:



“Heavy equipment” processors

Zen processors*




Advanced Vector eXtensions 512 – AVX-512

- 32 x 512-bit registers (Named: **ZMM0 – ZMM31**)
- Consists of multiple extensions
- Not meant to be supported by all processors implementing them.
- Only the core extension **AVX-512F (AVX-512 Foundation)** is required by all implementations.



AVX Family



- ✓ Each version incorporate new instructions.
- ✓ Backward compatibility is guaranteed.
- ✓ Each core can execute AVX instructions independently.
- ✓ But... there is a consideration...



AVX-512 “Problem”...

- ☐ AVX-512
- ☐ AVX-512F
- ☐ AVX-512BW
- ☐ AVX-512CD
- ☐ AVX-512DQ
- ☐ AVX-512ER
- ☐ AVX-512IFMA52
- ☐ AVX-512PF
- ☐ AVX-512VL
- ☐ AVX-512VPOPCNTDQ
- ☐ AVX-512_4FMAPS
- ☐ AVX-512_4VNNIW
- ☐ AVX-512_BF16
- ☐ AVX-512_BITALG
- ☐ AVX-512_VBMI
- ☐ AVX-512_VBMI2
- ☐ AVX-512_VNNI
- ☐ AVX-512_VP2INTERSECT

AVX-512? AVX-512-F?

Intel is introducing several variations of the new AVX-512 instruction support, based on the hardware and market segmentation. Ultimately there is one underlying set of AVX-512 instructions supported, and the different variants add different instructions.

AVX-512-F: F for Foundation

AVX-512-BW: Support for 512-bit Word support

AVX-512-CD: Conflict Detect (loop vectorization with possible conflicts)

AVX-512-DQ: More instructions for double/quad math operations

AVX-512-ER: Exponential and Reciprocal

AVX-512-IFMA: Integer Fused Multiply Add with 52-bit precision

AVX-512-PF: Prefetch Instructions

AVX-512-VBMI: Vector Byte Manipulation Instructions

AVX-512-VL: Foundation plus <512-bit vector length support

AVX-512-4VNNIW: Vector Neural Network Instructions Word (variable precision)

AVX-512-4FMAPS: Fused Multiply Accumulation Packed Single precision



Performance Benefit

➤ AVX, AVX2, AVX-512... Is it really worth it?

- 512-bit wide vectors
- 32 operand registers
- 8 64b mask registers
- Embedded broadcast
- Embedded rounding

Microarchitecture	Instruction Set	SP FLOPs / cycle	DP FLOPs / cycle
Skylake	Intel® AVX-512 & FMA	64	32
Haswell / Broadwell	Intel AVX2 & FMA	32	16
Sandybridge	Intel AVX (256b)	16	8
Nehalem	SSE (128b)	8	4

Intel AVX-512 Instruction Types	
AVX-512-F	AVX-512 Foundation Instructions
AVX-512-VL	Vector Length Orthogonality : ability to operate on sub-512 vector sizes
AVX-512-BW	512-bit Byte/Word support
AVX-512-DQ	Additional D/Q/SP/DP instructions (converts, transcendental support, etc.)
AVX-512-CD	Conflict Detect : used in vectorizing loops with potential address conflicts

Performance Benefit

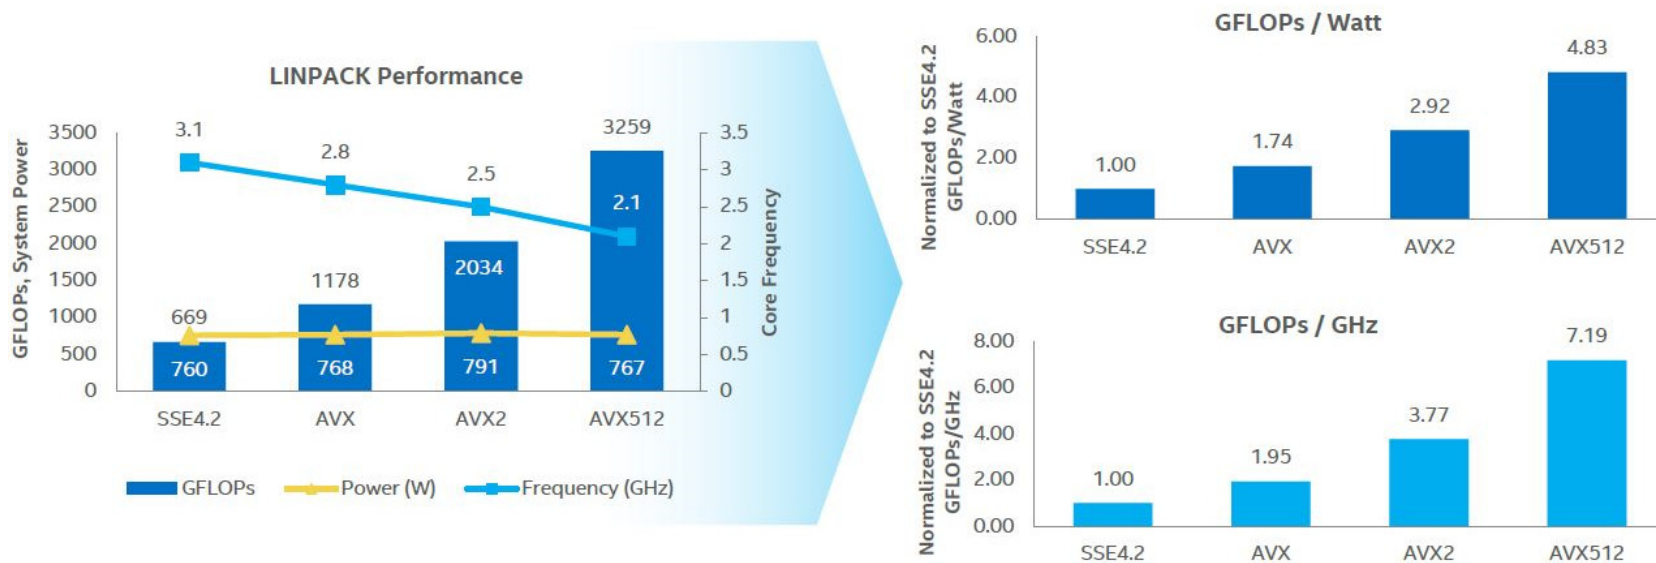
➤ AVX, AVX2, AVX-512... Is it really worth it?



<https://www.extremetech.com/computing/257730-intel-may-deploy-avx-512-upcoming-10nm-cannon-lake-cpus>

Performance Benefit

➤ AVX, AVX2, AVX-512... Is it really worth it?



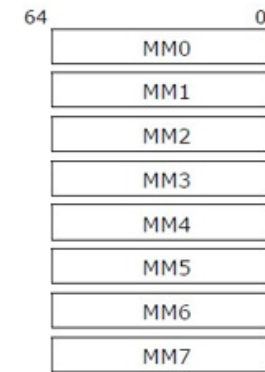
<https://www.extremetech.com/computing/257730-intel-may-deploy-avx-512-upcoming-10nm-cannon-lake-cpus>



MMX, SSE, AVX & AVX-512

Summary

- All SIMD Registers...



MMX

MMX, SSE, AVX & AVX-512

Summary

➤ All SIMD Registers...

64	0
MM0	
MM1	
MM2	
MM3	
MM4	
MM5	
MM6	
MM7	

MMX

SSE

127	0
XMM0	
XMM1	
XMM2	
XMM3	
XMM4	
XMM5	
XMM6	
XMM7	
XMM8	
XMM9	
XMM10	
XMM11	
XMM12	
XMM13	
XMM14	
XMM15	

MMX, SSE, AVX & AVX-512

Summary

➤ All SIMD Registers...

64	0
MM0	
MM1	
MM2	
MM3	
MM4	
MM5	
MM6	
MM7	

255	128	127	0
YMM0	XMM0		
YMM1	XMM1		
YMM2	XMM2		
YMM3	XMM3		
YMM4	XMM4		
YMM5	XMM5		
YMM6	XMM6		
YMM7	XMM7		
YMM8	XMM8		
YMM9	XMM9		
YMM10	XMM10		
YMM11	XMM11		
YMM12	XMM12		
YMM13	XMM13		
YMM14	XMM14		
YMM15	XMM15		

MMX



SSE

AVX

MMX, SSE, AVX & AVX-512


Summary

➤ All SIMD Registers...



511	256	255	128	127	0
ZMM0	YMM0	XMM0			
ZMM1	YMM1	XMM1			
ZMM2	YMM2	XMM2			
ZMM3	YMM3	XMM3			
ZMM4	YMM4	XMM4			
ZMM5	YMM5	XMM5			
ZMM6	YMM6	XMM6			
ZMM7	YMM7	XMM7			
ZMM8	YMM8	XMM8			
ZMM9	YMM9	XMM9			
ZMM10	YMM10	XMM10			
ZMM11	YMM11	XMM11			
ZMM12	YMM12	XMM12			
ZMM13	YMM13	XMM13			
ZMM14	YMM14	XMM14			
ZMM15	YMM15	XMM15			
ZMM16	YMM16	XMM16			
ZMM17	YMM17	XMM17			
ZMM18	YMM18	XMM18			
ZMM19	YMM19	XMM19			
ZMM20	YMM20	XMM20			
ZMM21	YMM21	XMM21			
ZMM22	YMM22	XMM22			
ZMM23	YMM23	XMM23			
ZMM24	YMM24	XMM24			
ZMM25	YMM25	XMM25			
ZMM26	YMM26	XMM26			
ZMM27	YMM27	XMM27			
ZMM28	YMM28	XMM28			
ZMM29	YMM29	XMM29			
ZMM30	YMM30	XMM30			
ZMM31	YMM31	XMM31			


You choose...



Is it possible to support more than one Extension Instruction Set?

It's possible:

- More code
- Bigger binaries
- Installation logic



CPU Internals...

CPU-Z

CPU | Caches | Mainboard | Memory | SPD | Graphics | Bench | About

Processor

Name	Intel Core i7 4700MQ		
Code Name	Haswell	Max TDP	47.0 W
Package	Socket 947 rPGA		
Technology	22 nm	Core Voltage	0.995 V

Specification Intel® Core™ i7-4700MQ CPU @ 2.40GHz

Family	6	Model	C	Stepping	3
Ext. Family	6	Ext. Model	3C	Revision	C0

Instructions MMX, SSE, SSE2, SSE3, SSSE3, SSE4.1, SSE4.2, EM64T, VT-x, AES, AVX, AVX2, FMA3

Clocks (Core #0)


Core Speed	3395.56 MHz
Multiplier	x 34.0 (8 - 34)
Bus Speed	99.87 MHz
Rated FSB	

Cache

L1 Data	4 x 32 KBytes	8-way
L1 Inst.	4 x 32 KBytes	8-way
Level 2	4 x 256 KBytes	8-way
Level 3	6 MBytes	12-way

Selection Socket #1 Cores 4 Threads 8


CPU-Z Ver. 1.90.0.x64 Tools Validate Close



CPU-Z

CPUID


<https://www.cpubid.com/>



Nickname: Axxxxxx - Nombre



Suggested Material



Introduction to x64 Assembly



AMD64 Technology





For next class...



What Happened to Cyrix Processors? | Nostalgia Nerd

<https://www.youtube.com/watch?v=iWGAdoMz1c0>



Thank you!

See you
next time!

